

# Boost Phoenix V3

Thomas Heller

Hartmut Kaiser



**Thomas Heller**

# Overview

- What's Phoenix?
  - Short History, Why a new version, Motivation
  - Phoenix as C++ in C++
  - Building a simple Asio echo server
- Boost Proto: The Phoenix' Workhorse
- Phoenix' Extension Mechanism
  - Building a parallel for construct
- Code as Data
- Compile Times

# What's Boost Phoenix?

# What's Phoenix? Short history

- Together with Joel de Guzman and Eric Niebler, Thomas implemented a new incarnation of the Boost.Phoenix library
- Started off as a Google Summer of Code project in 2010
- Passed Boost mini review early this year
- Will be part of Boost starting with V1.47, however it is in SVN already

# What's Phoenix? Why a new Version?

- Phoenix V2 was developed as a supporting library for Spirit
  - Features similar to Boost.Bind and Boost.Lambda
  - Hand-rolled Expression Templates (ET)
- Boost review in 2008
  - Unification of functional libraries with minimal disruption for users
  - Use Boost.Proto for unified placeholders and cross library integration
  - Use of C++11 features (rvalues, variadics)
- Improve and document extension mechanism

# What's Phoenix? Really?

- Enables functional programming techniques in C++
  - Higher order functions
  - Lambda (unnamed functions)
  - Currying (partial function application)
- C++ Embedded Domain Specific Language (EDSL) in C++
- Focus is more on usefulness and practicality than purity, elegance and strict adherence to FP principles

# Motivation

- Functional Programming style in the C++ standard library algorithms use function objects as “callbacks”

```
int init = 0;
std::accumulate (
    container.begin(), container.end(), init
    , std::plus<int>()
);
```

- Phoenix is the next evolutionary step

# Motivation

- Functional Programming style in the C++ standard library algorithms use function objects as “callbacks”

```
int init = 0;
std::accumulate (
    container.begin(), container.end(), init
    , arg1 + arg2
);
```

- Phoenix is the next evolutionary step

# Phoenix as C++ in C++

- Values and References
- Arguments
- Operators
- Statements
- Partial function application
- Construct, New, Delete, Casts
- Adapt arbitrary functions
- ...
- Everything is a function object!

# Phoenix as C++ in C++

- Values and References

```
cout << val(42)(); // 42
```

```
int i = 42;  
cout << ref(i)(); // 42
```

# Phoenix Function Objects

- Pseudo code for `val()`:

```
template <typename T>
struct val_impl {
    T value;
    val_impl(T t) : value(t) {}
    T operator()(...) const { return value; }
};
```

```
template <typename T> val_impl<T> val(T t)
{
    return val_impl<T>(t);
}
```

# Phoenix Function Objects

- Pseudo code for `ref()`:

```
template <typename T>
struct ref_impl {
    T& value;
    ref_impl(T& t) : value(t) {}
    T& operator()(...) const { return value; }
};
```

```
template <typename T> ref_impl<T> ref(T& t)
{
    return ref_impl<T>(t);
}
```

# Phoenix as C++ in C++

- Values and References
- Arguments

```
cout << arg1(42); // 42
```

```
const char* s = "Hello World!";  
cout << arg1(s); // Hello World!
```

- Alternative placeholder names available

```
cout << _1(i); // 42  
cout << _1(s); // Hello World!
```

# Phoenix Function Objects

- Pseudo code for arg1:

```
struct arg1_impl {  
    template <typename T1, ...>  
    T1 operator()(T1 t1, ...) const { return t1; }  
};
```

```
arg1_impl const arg1 = arg1_impl();
```

# Phoenix as C++ in C++

- Values and References
- Arguments
- Operators

```
cout << (arg1 * arg2)(2, 3);           // 6
```

```
int x = 3, z = 5;  
(ref(x) = arg1 + ref(z))(4);  
cout << x;                             // 9
```

```
(arg1 = arg2 + (3 * arg3))(ref(x), 3, 4)  
cout << x;                             // 15
```

# Phoenix Function Objects

- Pseudo code for +:

```
template <typename F1, typename F2> struct plus_impl {  
    F1 f1; F2 f2;  
    plus_impl(F1, f1, F2 f2) : f1(f1), f2(f2) {}  
    template <typename T1, ...>  
    T1 operator()(T1 t1, ...) const  
    { return f1(t1, ...) + f2(t1, ...); }  
};
```

```
template <typename F1, typename F2>  
plus_impl<F1, F2> operator+(F1 f1, F2 f2)  
{  
    return plus_impl<F1, F2>(f1, f2);  
}
```

# Phoenix as C++ in C++

- Values and References
- Arguments
- Operators
  - Unary prefix: `~`, `!`, `-`, `+`, `++`, `--`, `&` (reference), `*` (dereference)
  - Unary postfix: `++`, `--`
  - Binary: `=`, `[]`, `+=`, `-=`, `*=`, `/=`, `%=`, `&=`, `|=`, `^=`, `<<=`, `>>=`, `+`, `-`, `*`, `/`, `%`, `&`, `|`, `^`, `<<`, `>>`, `==`, `!=`, `<`, `>`, `<=`, `>=`, `&&`, `||`, `->*`
  - Ternary: `if_else(c, a, b)`

# Phoenix as C++ in C++

- Values and References
- Arguments
- Operators
- Statements

```
cout << (if_(arg1 > 5)[std::cout << arg1])(6); // 6
```

```
vector<int> v = { 4, 5, 6, 7 };  
std::for_each(v.begin(), v.end(),  
    if_(arg1 > 5)[  
        cout << arg1 << ", "  
    ]); // 6, 7
```

# Phoenix as C++ in C++

- Values and References
- Arguments
- Operators
- Statements
- Partial function application

```
void foo (int x, int y)
{
    std::cout << x+y << std::endl;
}
int i = 4;
bind(&foo, arg1, 3)(i); // 7
```

# Phoenix as C++ in C++

- Values and References
- Arguments
- Operators
- Statements
- Partial function application
- **Construct, New, Delete, Casts**
  - `construct<std::string>(arg1, arg2)`
  - `new_<std::string>(arg1, arg2)`
  - `delete_(arg1)`
  - `static_cast_<int*>(arg1)`

# Phoenix as C++ in C++

- Values and References
- Arguments
- Operators
- Statements
- Partial function application
- Construct, New, Delete, Casts
- Adapt arbitrary functions
  - Factorial

```
cout << fact(4)(); // 24
```

```
vector<double> v = { 1, 2, 3, 4 };  
for_each(v.begin(), v.end(), cout << fact(_1)); // 1 2 6 24
```

# Phoenix Function Function Objects

- Factorial code:

```
struct fact_impl {
    template <typename Sig>
    struct result;

    template <typename This, typename Arg>
    struct result<This(Arg const &)>
    { typedef Arg type; }

    template <typename Arg>
    Arg operator()(Arg const & n) const
    { return (n <= 1) ? 1 : n * (*this)(n-1); }
};
function<fact_impl> const fact = fact_impl();
```

# Phoenix vs. C++11 Lambdas

## C++11 Lambdas

- ✓ Built-in language feature
- ✓ No significant compile time hit
- Constructs monomorphic function objects
- ❖ Expressions need to be wrapped into lambda syntax

## Phoenix

- ✓ Library
- ✓ Expression are placed directly into code
- ✓ Constructs polymorphic function objects
- ✓ Constructs variadic function objects
- ❖ Significant compile time hit

# More Examples

# More Examples

```
// simple function object invocation
template <typename F>
void print(F f) { std::cout << f() << std::endl; }

print(val(3)); // 3
print(val("Hello World")); // Hello World


// find the first odd number in a vector
std::vector<int> c = { 2, 10, 4, 5, 1, 6, 8, 3, 9, 7 };
auto it = std::find_if(c.begin(), c.end(), arg1 % 2 == 1);
if (it != c.end()) cout << *it << std::endl;
```

# Full Example

Minimal Asio Echo Server

# Full Example: Asio Echo Server

- Adapt existing code, for example ASIO
  - Writing a simple asynchronous echo server



# Full Example: Asio Echo Server

- Adapt existing code, for example ASIO

read:

```
BOOST_PHOENIX_ADAPT_FUNCTION(void, read, asio::async_read, 3)
```

write:

```
BOOST_PHOENIX_ADAPT_FUNCTION(void, write, asio::async_write, 3)
```

buffer:

```
BOOST_PHOENIX_ADAPT_FUNCTION(  
    asio::mutable_buffers_1, buffer, asio::buffer, 2)
```

# Full Example: Asio Echo Server

- Adapt existing code, for example ASIO

accept:

```
template <typename Acceptor, typename Socket, typename Handler>
void accept_impl(Acceptor& acceptor, Socket& socket,
    Handler const& handler)
{
    acceptor.async_accept(socket, handler);
}
```

```
BOOST_PHOENIX_ADAPT_FUNCTION(void, accept, accept_impl, 3)
```

# Full Example: Asio Echo Server

- Adapt existing code, for example ASIO
- Using these functions we construct handlers:
  - `read_handler`  
`boost::function<void(system::error_code&, size_t)>`
  - `write_handler`  
`boost::function<void(system::error_code&, size_t)>`
  - `accept_handler`  
`boost::function<void(system::error_code&)>`

# Full Example: Asio Echo Server

- Adapt existing code, for example ASIO
  - `accept_handler`, will be called by ASIO when connection was accepted (or refused)

```
boost::function<void(system::error_code&)>
    accept_handler =
        if_(!_1)[
            read(ref(socket), buffer(ref(buf), max_length),
                ref(read_handler))
        ]
        .else [
            bind(&asio::ip::tcp::socket::close, ref(socket)),
            accept(ref(acceptor), ref(socket), ref(accept_handler))
        ];
```

# Full Example: Asio Echo Server

- Adapt existing code, for example ASIO
  - `read_handler`: will be called by ASIO after `read_async`

```
boost::function<void(system::error_code&, size_t)>
read_handler =
    if_( !_1)[
        write(ref(socket), buffer(ref(buf), _2),
              ref(write_handler)
            )
    ]
    .else [
        bind(&asio::ip::tcp::socket::close, ref(socket)),
        accept(ref(acceptor), ref(socket), ref(accept_handler))
    ];
```

# Full Example: Asio Echo Server

- Adapt existing code, for example ASIO
  - `write_handler`: will be called by ASIO after `write_async`

```
boost::function<void(system::error_code&, size_t)>
write_handler =
    if_(!_1)[
        read(ref(socket), buffer(ref(buf), max_length),
            ref(read_handler)
        ]
    .else [
        bind(&asio::ip::tcp::socket::close, ref(socket)),
        accept(ref(acceptor), ref(socket), ref(accept_handler))
    ];
```

# Full Example: Asio Echo Server

- Adapt existing code, for example ASIO
  - Start server

```
asio::io_service io_service;  
asio::ip::tcp::acceptor acceptor(io_service,  
    asio::ip::tcp::endpoint("localhost:1234"));  
asio::ip::tcp::socket socket(io_service);  
  
acceptor.async_accept(socket, accept_handler);
```

# Boost Proto

The workhorse behind Phoenix V3

# Boost Proto: The Phoenix' Workhorse

- The facilities of proto are used to form the backend of phoenix
  - Creation of expression template (ET) classes and composition
  - Formulation of transformations on the created ET tree

# Boost Proto: The Phoenix' Workhorse

- Provides facilities to generate your abstract syntax tree (AST)
  - Representation of your expression in terms of `proto::expr<>`, a hierarchical type holding terms of the expression by reference
- Describe your EDSL in terms of grammar rules
  - Not all expressions are valid
  - These grammars describe valid expressions
- Based on these rules, perform actions on the AST
  - Every rule is associated with an action, which is executed for the node the rule matched

# Boost Proto: Phoenix' Workhorse

- The type generated from a Proto expression is an instantiation of the Proto expression class

```
_1 + _2
```

```
template <typename Lhs, typename Rhs>  
proto::expr<proto::tag::plus, Lhs, Rhs>  
operator+ (Lhs const & lhs, Rhs const & rhs)  
{  
    return proto::make_expr<proto::tag::plus>(lhs, rhs);  
}
```

**SIMPLIFIED!**

# Boost Proto: Phoenix' Workhorse

- This type can be interpreted as a tree. This tree is what we call the AST for our language.

`1 + 2 → proto::make_expr<proto::tag::plus>(1, 2)`


# Boost Proto: Phoenix' Workhorse

- This type can be interpreted as a tree. This tree is what we call the AST for our language.

```
val(1) + 2 → proto::make_expr<proto::tag::plus>(1, 2)
```

```
proto::expr<  
  proto::tag::plus, proto::list2<  
    proto::expr<proto::tag::terminal, int>,  
    proto::expr<proto::tag::terminal, int>  
  > >
```


```
plus(terminal(1), terminal(2))
```



# The Phoenix AST in Proto's world

- Every Phoenix construct can be seen as an AST node of our Phoenix EDSL
- By composing these we create a bigger AST

```
if_( _1 == 0 )  
[  
    std::cout << _2  
]
```



# The Phoenix AST in Proto's world

- By using Proto we are able to introspect and transform this AST in any way we like
- The (lazy) evaluation inside Phoenix can be seen as a transformation of this AST
- Default (predefined) evaluation in Phoenix corresponds to 'normal' operator semantics
- By defining your own nodes you can extend Phoenix in any way you wish
- By customizing the evaluation of certain (predefined) nodes you can change the overall scheme
- By transforming the tree before evaluation you can do additional tricks

# Phoenix' Extension Mechanism

Add New Constructs: Define your own Node

# Phoenix' Extension Mechanism

Add new constructs

- Let's start simple. Assume you have the following code:

```
#pragma omp parallel for
for (int i = 0; i < NUM; ++i)
    c[i] = a[i] + b[i];
```

- And you want to express exactly this with Phoenix. How would you do that?

# Phoenix' Extension Mechanism

## Add new constructs

- Anticipated syntax:

```
omp::for_(<init>, <cond>, <reinit>)  
[  
    <parallel work>  
]
```

- Phoenix expression takes 4 arguments
  - Each is a Phoenix expression on its own

```
omp::for_(<init>, <cond>, <reinit>)  
[  
    <parallel work>  
]
```

Boost Phoenix V3 5/16/2011

# Phoenix' Extension Mechanism

## Add new constructs

- Step 1: A function object:

```
namespace omp  
{  
    struct for_eval  
    {  
        typedef void result_type;  
        template <typename I, typename C, typename S, typename D,  
                typename Ctx>  
        void operator()(I init, C cond, S step, D do_, Ctx ctx) const  
        {  
            #pragma omp parallel for  
            for(eval(init, ctx); eval(cond, ctx); eval(step, ctx))  
                eval(do_, ctx);  
        }  
    };  
}
```

```
omp::for_(<init>, <cond>, <reinit>)
[
  <parallel work>
]
```

Boost Phoenix V3 5/16/2011

# Phoenix' Extension Mechanism

Add new constructs

- Step 2: Create the Phoenix expression type:

```
BOOST_PHOENIX_DEFINE_EXPRESSION(
  (omp)(for ) // define omp::for_
  , (boost::phoenix::meta_grammar) // Cond
  (boost::phoenix::meta_grammar) // Init
  (boost::phoenix::meta_grammar) // Step
  (boost::phoenix::meta_grammar) // Do
)
```

Four  
Arguments

- Defines:

```
omp::make_for_, omp::result_of::make_for_, omp::expression::for_
omp::rule::for_, omp::tag::for_, omp::functional::make_for_
```

```
omp::for_(<init>, <cond>, <reinit>)  
[  
    <parallel work>  
]
```

Boost Phoenix V3 5/16/2011

# Phoenix' Extension Mechanism

## Add new constructs

- Step 3: Create expression generator: `omp::for_`

```
namespace omp {  
    template <typename Init, typename Cond, typename Step>  
    struct for_gen {...};  
  
    template <typename Init, typename Cond, typename Step>  
    inline for_gen<Init, Cond, Step> const  
    for_(Init const& init, Cond const& cond, Step const& step)  
    {  
        return for_gen<Init, Cond, Step>(init, cond, step);  
    }  
}
```

```
omp::for_(<init>, <cond>, <reinit>)
[
  <parallel work>
]
```

Boost Phoenix V3 5/16/2011

# Phoenix' Extension Mechanism

## Add new constructs

- Step 3: Create expression generator

```
template <typename Init, typename Cond, typename Step>
struct for_gen {
  for_gen(Init const& init, Cond const& cond, Step const& step)
    : init(init), cond(cond), step(step) {}
};
```

```
template <typename Do>
typename omp::result_of::make_for_<Init, Cond, Step, Do>::type const
operator[](Do const& do_) const
{
  return omp::make_for_(init, cond, step, do_);
}

Init init; Cond cond; Step step;
};
```

Phoenix "Magic"

```
omp::for_(<init>, <cond>, <reinit>)  
[  
    <parallel work>  
]
```

Boost Phoenix V3 5/16/2011

# Phoenix' Extension Mechanism

## Add new constructs

- Step 4: Define how to evaluate the new expression

```
namespace boost { namespace phoenix  
{  
    template <>  
    struct default_actions : when<omp::rule::for_>  
        : phoenix::call<omp::for_eval>  
    {};  
}}
```

Use ...::when to  
associate a grammar rule  
with an action

```
omp::for_(<init>, <cond>, <reinit>)  
[  
    <parallel work>  
]
```

Boost Phoenix V3 5/16/2011

# Phoenix' Extension Mechanism

## Add new constructs

- Step 5: Use it

```
std::vector<int> a(NUM, 1), b(NUM, 2), c(NUM, 0);  
(  
    let(_a = begin(_1), _b = begin(_2), _c = begin(_3))  
    [  
        omp::for_(nothing, _a != end(_1), (++_a, ++_b, ++_c))  
        [  
            *_c = *_a + *_b           // executed in parallel!  
        ]  
        , std::cout << accumulate(_3, 0) << "\n"  
    ]  
)  
(a, b, c);
```

# Phoenix' Extension Mechanism

Customizing Existing Constructs: Changing the Evaluation Scheme

# Phoenix' Extension Mechanism

## Reusing for\_ - Changing the evaluation scheme

- Anticipated syntax:

```
omp::parallel(  
    <parallel work>  
    for_(<init>, <cond>, <reinit>) [ <body> ]  
)
```

- Any Phoenix expression inside `omp::parallel()` will be (optionally) parallelized
  - We show how to parallelize `for_`

```
omp::parallel(  
  <parallel work>  
  for_(<init>, <cond>, <reinit>) [ <body> ]  
)
```

Boost Phoenix V3 5/16/2011

# Phoenix' Extension Mechanism

## Reusing for\_ - Changing the evaluation scheme

- Step 1: Define how to evaluate the new expression

```
// Define new a new action  
namespace omp {  
  struct parallel_actions  
  {  
    template <typename Rule>  
    struct when : phoenix::default_actions::when<Rule>  
    {};  
  };  
  
  // only change what we are interested in:  
  template <>  
  struct parallel_actions::when<phoenix::rule::for_>  
  : phoenix::call<omp::for_eval>  
  {};  
}
```

```
omp::parallel(
  <parallel work>
  for_(<init>, <cond>, <reinit>) [ <body> ]
)
```

Boost Phoenix V3 5/16/2011

# Phoenix' Extension Mechanism

## Reusing for\_ - Changing the evaluation scheme

- Step 2: Change the evaluation scheme on the fly

```
// omp::make_parallel, omp::rule::parallel, etc.
BOOST_PHOENIX_DEFINE_EXPRESSION(
  (omp)(parallel), (phoenix::meta_grammar))

namespace boost { namespace phoenix
{
  template <>
  struct default_actions::when<omp::rule::parallel>
  : proto::call<phoenix::evaluator(
    proto::_child0,
    phoenix::functional::context(
      phoenix::_env, omp::parallel_actions()),
    phoenix::unused)>
  {};
}}
```

```
omp::parallel(  
    <parallel work>  
    for_(<init>, <cond>, <reinit>) [ <body> ]  
)
```

Boost Phoenix V3 5/16/2011

# Phoenix' Extension Mechanism

## Reusing for\_ - Changing the evaluation scheme

- Step 3: Create expression generator

```
namespace omp  
{  
    template <typename Expr>  
    typename omp::result_of::make_parallel<Expr>::type const  
    parallel(Expr const& expr)  
    {  
        return omp::make_parallel(expr);  
    }  
}
```



One Argument

```
parallel(
  <parallel work>
  for_(<init>, <cond>, <reinit>) [ <body> ]
)
```

Boost Phoenix V3 5/16/2011

# Phoenix' Extension Mechanism

## Reusing for\_ - Changing the evaluation scheme

- Step 4: Use it

```
std::vector<int> a(NUM, 1), b(NUM, 2), c(NUM, 0);
(
  let(_a = begin(_1), _b = begin(_2), _c = begin(_3))
  [
    omp::parallel( // now being executed in parallel!
      for_(nothing, _a != end(_1), (++_a, ++_b, ++_c))
      [
        *_c = *_a + *_b
      ]
    )
  ], std::cout << accumulate(_3, 0) << "\n"
)
](a, b, c);
```

# Phoenix' Extension Mechanism

Code as Data

# Code as Data

“Any sufficiently complicated C or Fortran program contains an ad hoc, informally-specified, bug-ridden, slow implementation of half of Common Lisp”

Greenspun's Tenth Rule of Programming

# Code as Data

- We now have both,
  - the possibility of implementing custom expressions and
  - the ability to transform these expressions,
- We gained a powerful implementation of scheme like AST macro capabilities at compile time

# Code as Data

## Example – Using proto to transform the AST

- Inverting arithmetic expressions

```
struct invert_actions
{
    // By default, just return the expression itself
    template <typename Rule>
    struct when : proto::_
    {};
};
```

# Code as Data

## Example – Using proto to transform the AST

- Inverting arithmetic expressions

```
// When there is a plus, make a minus out of it
template <>
struct invert_actions::when<phoenix::rule::plus>
    : proto::call<
        phoenix::functional::make_minus(
            phoenix::evaluator(
                proto::_left, phoenix::_context, phoenix::unused),
            phoenix::evaluator(
                proto::_right, phoenix::_context, phoenix::unused)
        )
    >
{};
```

# Code as Data

## Example – Using proto to transform the AST

- Inverting arithmetic expressions

```
// When there is a minus, make a plus out of it
template <>
struct invert_actions::when<phoenix::rule::minus>
    : proto::call<
        phoenix::functional::make_plus(
            phoenix::evaluator(
                proto::_left, phoenix::_context, phoenix::unused),
            phoenix::evaluator(
                proto::_right, phoenix::_context, phoenix::unused)
        )
    >
{};
```

# Code as Data

## Example – Using proto to transform the AST

- Creating the generator


```
template <typename Expr>
typename phoenix::result_of::eval<
    Expr const&,
    phoenix::result_of::make_context<
        phoenix::make_env<>, invert_actions>::type
>::type
invert(Expr const& expr)
{
    return phoenix::eval(
        expr,
        phoenix::make_context(phoenix::make_env(), invert_actions())
    );
}
```

```
(_1 * invert(_2 - _3))(2, 3, 4);
```

```
// 14
```

# Phoenix Compile Times

# Phoenix Compile Times



- Effect of partial preprocessing of Phoenix Headers on compile times
  - No PP: without partially preprocessed headers
  - Using PP: when using partially preprocessed headers

# Phoenix Compile Times

- T1: `bind_member_function_tests.cpp`

<b>PHOENIX_LIMIT = 10</b>	<b>T1 (gcc)</b>	<b>T1 (VC2010)</b>
Phoenix V2	2.9 [s]	2.58 [s]
Phoenix V3, no PP	3.4 [s]	4.05 [s]
Phoenix V3, no PP, Proto PP	2.6 [s]	3.43 [s]
Phoenix V3, full PP	2.3 [s]	2.96 [s]

# Conclusions

# Imagine the unimaginable

- We modeled C++ inside C++
- With the help of Proto, we created a powerful compiler toolkit
- Enabling the creation of new technologies:
  - Multi stage programming, completely done in C++
  - Optimize code based on the high level information of the AST
  - Change the evaluation of a Phoenix expression to whatever you like