

Object-Relational Mapping with ODB and Boost

Boris Kolpackov

Code Synthesis

v1.2, May 2011

**CODE
SYNTHESIS**

ORM for C++

No

- No manual parameter binding
- No manual result set extraction
- No hand-written mapping or registration code
- Not a framework

ORM for C++

Yes

- Automatic generation of mapping code from C++ classes
- Handle any standard C++ code
- Object-oriented database API
- Statically-typed, C++-integrated query language
- Database portability
- Flexible and customizable

GCC Plugin Support

- Available since GCC 4.5.0/April 2010
- Current releases are GCC 4.5.3 and 4.6.0

GCC Plugin Support

- Available since GCC 4.5.0/April 2010
- Current releases are GCC 4.5.3 and 4.6.0

- Dynamic loading
- Hook into compilation pipeline *anywhere*
- ...starting from compiler startup
- ...ending with assembler output

Platforms

ODB can be used with any modern C++ compiler

- GNU/Linux with GCC
- Mac OS X with GCC
- Solaris (x86 and SPARC) with Sun Studio C++
- Windows with GCC (MinGW) and VC++ 2008 and 2010

Databases

Cross-database

Supported

MySQL
SQLite

Coming Soon

PostgreSQL
Oracle
MS SQL

Performance

High-performance and low overhead

- Prepared statements
- Caching of connections, statements, and buffers
- Low-level native database APIs
- Zero per-object memory overhead

Performance

High-performance and low overhead

- Prepared statements
- Caching of connections, statements, and buffers
- Low-level native database APIs
- Zero per-object memory overhead

Loading performance

- MySQL — 21,000 objects per second — 48 μ s per object
- SQLite — 143,000 object per second — 7 μ s per object

License

Dual-licensed

- GPL + commercial license
- Can be used without restrictions within your organization

Declaring a Persistent Class

```
enum status {open, confirmed, closed};

class bug
{
public:
    bug (const std::string& summary,
         const std::string& description);
    ...

private:
    unsigned int id_;

    status status_;
    std::string summary_;
    std::string description_;
};
```

Declaring a Persistent Class

```
enum status {open, confirmed, closed};
```

```
#pragma db object
```

```
class bug
```

```
{
```

```
    ...
```

```
private:
```

```
    friend class odb::access;
```

```
    bug () {}
```

```
#pragma db id auto
```

```
unsigned int id_;
```

```
status status_;
```

```
std::string summary_;
```

```
std::string description_;
```

```
};
```

Declaring a Persistent Class

```
enum status {open, confirmed, closed};
```

```
#pragma db object
```

```
class bug
```

```
{
```

```
    ...
```

```
private:
```

```
    friend class odb::access;
```

```
    bug () {}
```

```
#pragma db id auto
```

```
unsigned int id_;
```

```
status status_;
```

```
std::string summary_;
```

```
std::string description_;
```

```
};
```

Declaring a Persistent Class

```
enum status {open, confirmed, closed};
```

```
#pragma db object
```

```
class bug
```

```
{
```

```
    ...
```

```
private:
```

```
    friend class odb::access;
```

```
    bug () {}
```

```
#pragma db id auto
```

```
unsigned int id_;
```

```
    status status_;
```

```
    std::string summary_;
```

```
    std::string description_;
```

```
};
```

Declaring a Persistent Class

```
enum status {open, confirmed, closed};
```

```
#pragma db object
```

```
class bug
```

```
{
```

```
    ...
```

```
private:
```

```
    friend class odb::access;
```

```
    bug () {}
```

```
#pragma db id auto
```

```
unsigned int id_;
```

```
    status status_;
```

```
    std::string summary_;
```

```
    std::string description_;
```

```
};
```

Declaring a Persistent Class

```
enum status {open, confirmed, closed};
```

```
#pragma db object
```

```
class bug
```

```
{
```

```
    ...
```

```
private:
```

```
    friend class odb::access;
```

```
    bug () {}
```

```
#pragma db id auto
```

```
unsigned int id_;
```

```
status status_;
```

```
std::string summary_;
```

```
std::string description_;
```

```
};
```


Declaring a Persistent Class

```
enum status {open, confirmed, closed};
```

```
#pragma db object
```

```
class bug
```

```
{
```

```
    ...
```

```
private:
```

```
    friend class odb::access;
```

```
    bug () {}
```

```
#pragma db id auto
```

```
unsigned int id_;
```

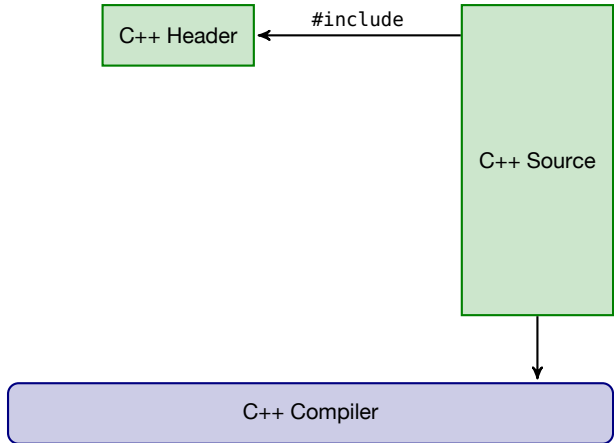
```
    status status_;
```

```
    std::string summary_;
```

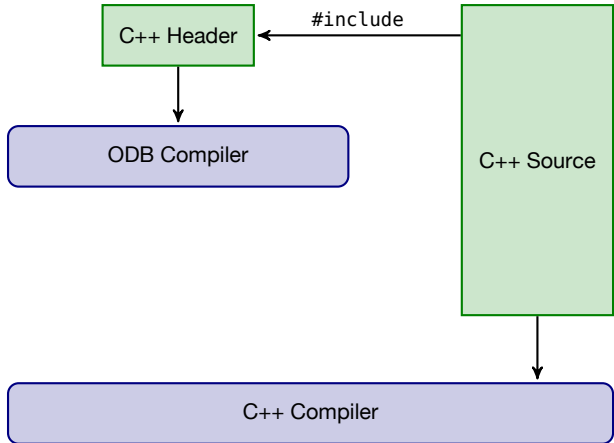
```
    std::string description_;
```

```
};
```

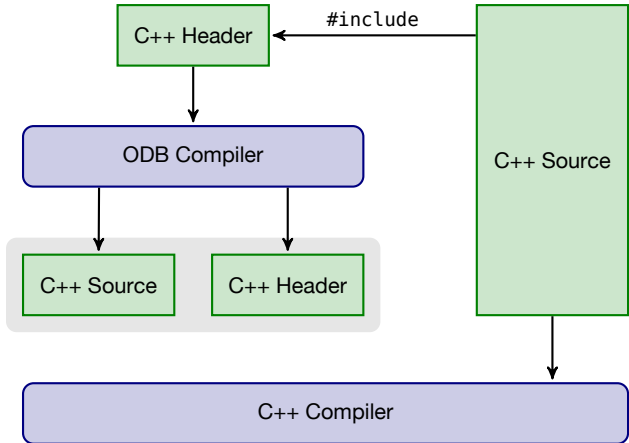
Workflow



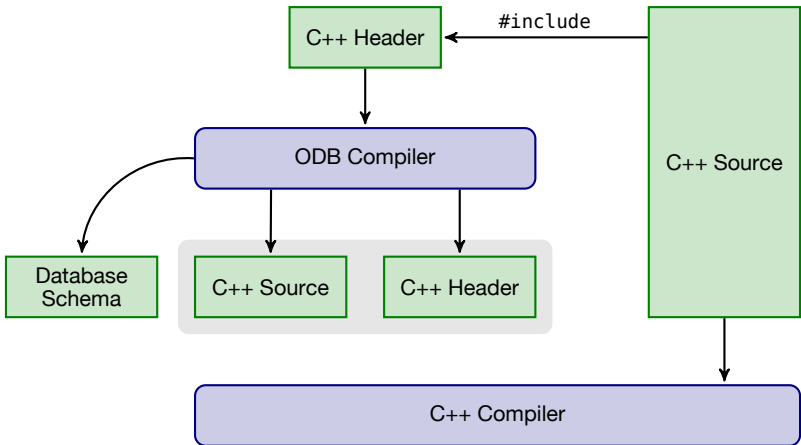
Workflow



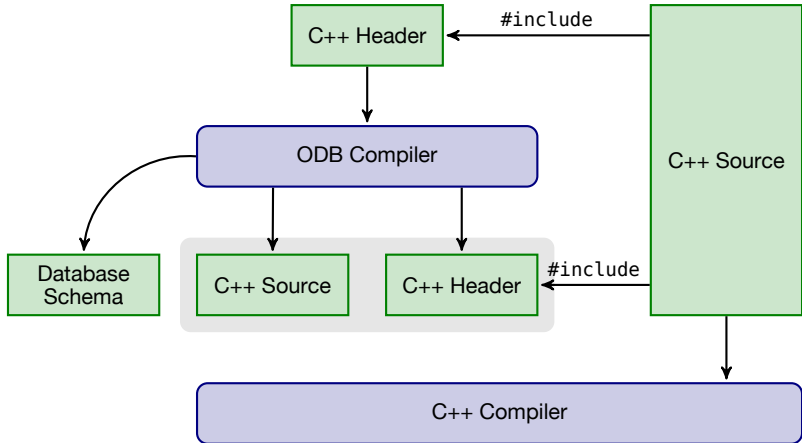
Workflow



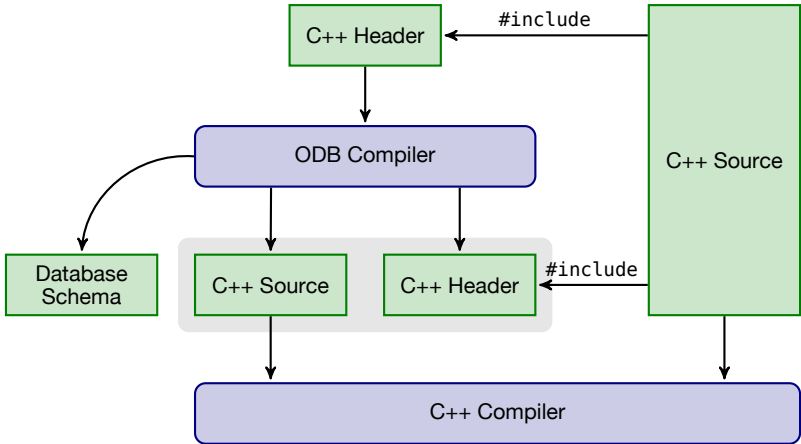
Workflow



Workflow



Workflow



ODB Compiler

```
odb --database mysql --generate-schema bug.hxx
```


ODB Compiler

```
odb --database mysql --generate-schema bug.hxx
```

```
odb -I/opt/boost -d mysql -s bug.hxx
```

ODB Compiler

```
odb --database mysql --generate-schema bug.hxx
```

```
odb -I/opt/boost -d mysql -s bug.hxx
```

```
CREATE TABLE bug (  
  id INT UNSIGNED NOT NULL PRIMARY KEY AUTO_INCREMENT,  
  status ENUM('open', 'confirmed', 'closed') NOT NULL,  
  summary TEXT NOT NULL,  
  description TEXT NOT NULL);
```

Database

```
#include <odb/database.hxx>
#include <odb/mysql/database.hxx>

shared_ptr<odb::database> db (
    new odb::mysql::database ("bugtracker", // user
                             "secret",    // password
                             "bugs"));    // database
```

Database

```
#include <odb/database.hxx>
#include <odb/mysql/database.hxx>
```

```
shared_ptr<odb::database> db (
    new odb::mysql::database ("bugtracker", // user
                              "secret",    // password
                              "bugs"));    // database
```

```
#include <odb/sqlite/database.hxx>
```

```
shared_ptr<odb::database> db (
    new odb::sqlite::database ("bugs.db")); // database
```

Making an Object Persistent

```
bug b ("Support for PostgreSQL",  
      "ODB does not yet support PostgreSQL.");  
  
transaction t (db->begin ());  
  
unsigned int id = db->persist (b);  
  
t.commit ();
```

Making an Object Persistent

```
bug b ("Support for PostgreSQL",  
      "ODB does not yet support PostgreSQL.");  
  
transaction t (db->begin ());  
  
unsigned int id = db->persist (b);  
  
t.commit ();
```

```
=> INSERT INTO bug (  
    id,  
    status,  
    summary,  
    description)  
VALUES (?, ?, ?, ?)
```

Loading a Persistent Object

```
transaction t (db->begin ());  
  
shared_ptr<bug> b (db->load<bug> (id));  
  
// bug b;  
// db->load (id, b);  
  
t.commit ();
```

Loading a Persistent Object

```
transaction t (db->begin ());  
  
shared_ptr<bug> b (db->load<bug> (id));  
  
// bug b;  
// db->load (id, b);  
  
t.commit ();
```

```
=> SELECT  
    status,  
    summary,  
    description  
FROM bug WHERE id = ?
```


Updating a Persistent Object

```
transaction t (db->begin ());  
  
shared_ptr<bug> b (db->load<bug> (id));  
b->confirm ();  
db->update (b);  
  
t.commit ();
```

Updating a Persistent Object

```
transaction t (db->begin ());  
  
shared_ptr<bug> b (db->load<bug> (id));  
b->confirm ();  
db->update (b);  
  
t.commit ();
```

```
=> UPDATE bug SET  
    status = ?,  
    summary = ?,  
    description = ?  
WHERE id = ?
```

Querying the Database

```
typedef odb::query<bug> query;
typedef odb::result<bug> result;

transaction t (db->begin ());

result r (db->query<bug> (query::status == open));

for (result::iterator i (r.begin ()); i != r.end (); ++i)
    cout << i->id () << " " << i->summary () << endl;

t.commit ();
```

Querying the Database

```
typedef odb::query<bug> query;
typedef odb::result<bug> result;

transaction t (db->begin ());

result r (db->query<bug> (query::status == open));

for (result::iterator i (r.begin ()); i != r.end (); ++i)
    cout << i->id () << " " << i->summary () << endl;

t.commit ();
```

Querying the Database

```
typedef odb::query<bug> query;
typedef odb::result<bug> result;

transaction t (db->begin ());

result r (db->query<bug> (query::status == open));

for (result::iterator i (r.begin ()); i != r.end (); ++i)
    cout << i->id () << " " << i->summary () << endl;

t.commit ();
```

```
=> SELECT
      id
      status,
      summary,
      description
FROM bug WHERE status = ?
```

Querying the Database

```
db->query<bug> (query::status == open ||
               query::status == confirmed);
```

```
status s;
query q (query::status == query::_ref (s));
```

```
s = open;
db->query<bug> (q); // status == open
```

```
s = closed;
db->query<bug> (q); // status == closed
```

```
db->query<bug> ("status = " + query::_val (open));
```

```
db->query<bug> ("stats = " + query::_val (123));
```

Querying the Database

```
db->query<bug> (query::status == open ||  
               query::status == confirmed);
```

```
status s;  
query q (query::status == query::_ref (s));
```

```
s = open;  
db->query<bug> (q); // status == open
```

```
s = closed;  
db->query<bug> (q); // status == closed
```

```
db->query<bug> ("status = " + query::_val (open));
```

```
db->query<bug> ("stats = " + query::_val (123));
```

Querying the Database

```
db->query<bug> (query::status == open ||
               query::status == confirmed);
```

```
status s;
query q (query::status == query::_ref (s));
```

```
s = open;
db->query<bug> (q); // status == open
```

```
s = closed;
db->query<bug> (q); // status == closed
```

```
db->query<bug> ("status = " + query::_val (open));
```

```
db->query<bug> ("stats = " + query::_val (123));
```


Querying the Database

```
db->query<bug> (query::status == open ||  
               query::status == confirmed);
```

```
status s;  
query q (query::status == query::_ref (s));
```

```
s = open;  
db->query<bug> (q); // status == open
```

```
s = closed;  
db->query<bug> (q); // status == closed
```

```
db->query<bug> ("status = " + query::_val (open));
```

```
db->query<bug> ("stats = " + query::_val (123));
```

Querying the Database

```
db->query<bug> (query::status == open ||  
               query::status == confirmed);
```

```
status s;  
query q (query::status == query::_ref (s));
```

```
s = open;  
db->query<bug> (q); // status == open
```

```
s = closed;  
db->query<bug> (q); // status == closed
```

```
db->query<bug> ("status = " + query::_val (open));
```

```
db->query<bug> ("stats = " + query::_val (123));
```

Deleting a Persistent Object

```
transaction t (db->begin ());  
db->erase<bug> (id);  
t.commit ();
```

Deleting a Persistent Object

```
transaction t (db->begin ());
```

```
db->erase<bug> (id);
```

```
t.commit ();
```

=> **DELETE FROM** bug **WHERE** id = ?

Profiles

- Generic integration mechanism
- Covers smart pointers, containers, and value types
- ODB includes profiles for Boost and Qt
- You can write your own

```
odb -d mysql -p boost bug.hxx
```

Boost Profile

- `shared_ptr/weak_ptr` + their lazy variants
- unordered containers library
- `date_time` library

Adding Creation and Modification Dates

```
#pragma db object
class bug
{
    ...

    #pragma db id auto
    unsigned int id_;

    status status_;
    std::string summary_;
    std::string description_;

    boost::posix_time::ptime created_;
    boost::posix_time::ptime updated_;
};
```

Containers

- Standard containers: vector, list, set, map, etc
- Profiles provide additional containers
- Easy to support custom containers

Adding Comments and Tags

```
#pragma db object
class bug
{
    ...

    #pragma db id auto
    unsigned int id_;

    status status_;
    std::string summary_;
    std::string description_;

    boost::posix_time::ptime created_;
    boost::posix_time::ptime updated_;

    std::vector<std::string> comments_;
    boost::unordered_set<std::string> tags_;
};
```

Composite Value Types

- Class or struct type
- Mapped to more than one database column
- Can contain composite values, containers, and pointers to objects

Extending Comments

```
#pragma db value
class comment
{
    ...

    std::string text_;
    boost::posix_time::ptime created_;
};
```

```
#pragma db object
class bug
{
    ...

    std::vector<comment> comments_;
};
```

Relationships

- Relationships are represented as pointers to objects
- Standard pointers: raw, auto_ptr, tr1::shared_ptr
- Profiles provide additional pointers
- Easy to support custom smart pointers

Adding User Object

```
#pragma db object
class user
{
    ...

    #pragma db id
    std::string email_;

    std::string first_;
    std::string last_;
};
```

Adding Reporter

```
#pragma db object  
class bug  
{  
    ...  
    shared_ptr<user> reporter_  
};
```

Adding Reporter

```
#pragma db object  
class bug  
{  
    ...  
    shared_ptr<user> reporter_  
};
```

Example of a *unidirectional to-one* relationship.

Adding Bug List

```
#pragma db object
class user
{
    ...

    #pragma db id
    std::string email_;

    std::string first_;
    std::string last_;

    #pragma db inverse(reporter_)
    std::vector<shared_ptr<bug> > reported_bugs_;
};
```


Adding Bug List

```
#pragma db object
class user
{
    ...

    #pragma db id
    std::string email_;

    std::string first_;
    std::string last_;

    #pragma db inverse(reporter_)
    std::vector<shared_ptr<bug> > reported_bugs_;
};
```

Example of a *bidirectional many-to-one* relationship.

Adding Bug List

```
#pragma db object
class user
{
    ...

    #pragma db id
    std::string email_;

    std::string first_;
    std::string last_;

    #pragma db inverse(reporter_)
    std::vector<shared_ptr<bug> > reported_bugs_;
};
```

Example of a *bidirectional many-to-one* relationship.

Inverse Side of a Relationship

```
CREATE TABLE user (  
  email VARCHAR (255) NOT NULL PRIMARY KEY,  
  ...);
```

```
CREATE TABLE user_reported_bugs (  
  object_id VARCHAR (255) NOT NULL,  
  index BIGINT UNSIGNED NOT NULL,  
  value INT UNSIGNED REFERENCES bug (id));
```

```
CREATE TABLE bug (  
  id INT UNSIGNED NOT NULL PRIMARY KEY AUTO_INCREMENT,  
  ...  
  reporter VARCHAR (255) REFERENCES user (email));
```

Inverse Side of a Relationship

```
CREATE TABLE user (  
  email VARCHAR (255) NOT NULL PRIMARY KEY,  
  ...);
```

```
CREATE TABLE bug (  
  id INT UNSIGNED NOT NULL PRIMARY KEY AUTO_INCREMENT,  
  ...  
  reporter VARCHAR (255) REFERENCES user (email));
```

We Have a Problem

```
#pragma db object
class user
{
    ...

    #pragma db inverse(reporter_)
    std::vector<shared_ptr<bug> > reported_bugs_;
};
```

```
#pragma db object
class bug
{
    ...

    shared_ptr<user> reporter_;
};
```

We Have a Problem

Actually, we have two:

1. Ownership cycle between user and bug
2. Eager loading of the bug list in user

We Have a Problem

Actually, we have two:

1. Ownership cycle between user and bug
2. Eager loading of the bug list in user

```
#pragma db object
```

```
class user
```

```
{
```

```
...
```

```
#pragma db inverse(reporter_)
```

```
std::vector<weak_ptr<bug> > reported_bugs_;
```

```
};
```

Lazy Pointers

- Finer grained control over relationship loading
- Every supported pointer has a corresponding lazy version

Lazy Pointers

- Finer grained control over relationship loading
- Every supported pointer has a corresponding lazy version

```
#pragma db object
class user
{
    ...

    #pragma db inverse(reporter_)
    std::vector<lazy_weak_ptr<bug> > reported_bugs_;
};
```

```
lazy_weak_ptr<bug> lwp = ...
shared_ptr<bug> b (lwp.load ()); // Load and lock.
```

Generated or Custom Schema

- Database schema can be automatically generated
- Or we can map persistent classes to a custom schema

Generated Schema

- Standalone SQL file
- Embedded into the generated C++ code

Generated Schema

- Standalone SQL file
- Embedded into the generated C++ code

```
shared_ptr<odb::database> db (  
    new odb::mysql::database (...));  
  
transaction t (db->begin ());  
schema_catalog::create_schema (*db);  
t.commit ();
```

Custom Schema

- Map classes to tables
- Map members to columns
- Map C++ types to database types

Legacy Bug Database

```
CREATE TABLE Bugs (  
  BugId INT UNSIGNED NOT NULL PRIMARY KEY AUTO_INCREMENT,  
  BugStatus INT NOT NULL,  
  Summary VARCHAR (128) NOT NULL,  
  Descr VARCHAR (1024) NOT NULL)
```

Legacy Bug Database

```
#pragma db object table("Bugs")
class bug
{
    ...

    #pragma db id auto column("BugId")
    unsigned int id_;

    #pragma db column("BugStatus") type("INT")
    status status_;

    #pragma db column("Summary") type("VARCHAR (128)")
    std::string summary_;

    #pragma db column("Descr") type("VARCHAR (1024)")
    std::string description_;
};
```

Mapping Placement

```
class bug
{
    ...

    unsigned int id_;
    status status_;
    std::string summary_;
    std::string description_;
};

#pragma db object(bug) table("Bugs")
#pragma db member(bug::id_) id auto column("BugId")
...
```


Mapping Placement

```
// bug.hxx
class bug
{
    ...

    unsigned int id_;
    status status_;
    std::string summary_;
    std::string description_;
};

#include "bug-mapping.hxx"
// bug-mapping.hxx
#pragma db object(bug) table("Bugs")
#pragma db member(bug::id_) id auto column("BugId")
...
```

Mapping Placement

```
// bug.hxx
```

```
class bug
{
    ...

    unsigned int id_;
    status status_;
    std::string summary_;
    std::string description_;
};
```

```
// bug-mapping.hxx
```

```
#pragma db object(bug) table("Bugs")
#pragma db member(bug::id_) id auto column("BugId")
...
```

```
odb --odb-epilogue-file bug-mapping.hxx bug.hxx
```

Summary

ODB can rival Hibernate for Java and EF/NHibernate for C#.

- Open-source
- Cross-platform
- Cross-database
- Well documented

Future

- Support more libraries from Boost
- Support more databases
- SQL to C++ compiler
- Optimistic concurrency
- Database schema evolution and versioning

Resources

- [ODB home page](#)
 - www.codesynthesis.com/products/odb/
- [ODB manual](#)
 - www.codesynthesis.com/products/odb/doc/manual.xhtml
- [ODB mailing lists](#)
 - www.codesynthesis.com/products/odb/ mailing-lists.xhtml
- [A Sense of Design](#)
 - www.codesynthesis.com/~boris/blog/

Questions

