

Practical SIMD acceleration with Boost.SIMD

Mathias Gaunard Joël Falcou Jean-Thierry Lapresté

MetaScale


Boostcon 2011

Context

From NT^2 to Boost.SIMD

- Last year, we presented NT^2 , a MATLAB-like Proto-based library for high-performance numerical computation
- Boost.SIMD is the extraction of the SIMD subcomponent of the library
- A GSoC project is scheduled this summer to help make it ready for review
- This talk is here to present what's inside the proposal.

What's SIMD?



Principles

- Single Instruction, Multiple Data
- Operations applied on $N \times T$ elements within a single register
- Up to N times faster than regular ALU/FPU

Why is SIMD abstraction needed?

x86 family

- MMX 64-bit float, double
- SSE 128-bit float
- SSE2 128-bit int8, int16, int32, int64, double
- SSE3
- SSSE3
- SSE4a (AMD only)
- SSE4.1
- SSE4.2
- AVX 256-bit float, double
- FMA4 (AMD only)
- XOP (AMD only)
- FMA3

PowerPC family

- AltiVec 128-bit int8, int16, int32, int64, float
- Cell SPU 128-bit int8, int16, int32, int64, float, double

ARM family

- VFP 64-bit float, double
- NEON 64-bit and 128-bit float, int8, int16, int32, int64

Why not let the compiler do it?

Compilers are only so smart

- Automatic vectorization can only happen if:
 - Memory is well agenced
 - Code is inherently vectorizable
- Compiled functions are not vectorized (I look at you libm !)
- Compilers don't always have enough static information to know what they can vectorize
- Designing for vectorization is a human process

Conclusion

- Declaring SIMD parallelism explicitly is the best way to get your code vectorized
- To be demonstrated by this presentation

Talk Layout

- 1 Introduction
- 2 Interface**
- 3 SIMD and STL
- 4 SIMD Specific Idioms
- 5 Conclusion

Writing it by hand

Doing $a * b + c$ with vectors of 32-bit integers : SSE

```
__m128i a, b, c, result;  
result = _mm_mul_epu32(a, _mm_add_epu32(b, c));
```

Doing $a * b + c$ with vectors of 32-bit integers : AltiVec

```
__vector int a, b, c, result;  
result = vec_cts(vec_madd( vec_ctf(a,0)  
                          , vec_ctf(b,0)  
                          , vec_ctf(c,0)  
                          )  
                ,0);
```

The pack abstraction

```
simd::pack<T>
```

`pack<T, N>` SIMD register that packs N elements of type T
`pack<T>` automatically finds best N available

- Behaves just like T except operations yield a pack of T and not a T.

Constraints

- T must be a fundamental arithmetic type, i.e. `(un)signed char`, `(unsigned) short`, `(unsigned) int`, `(unsigned) long`, `(unsigned) long long`, `float` or `double` – not `bool`.
- N must be a power of 2.

pack API

Operators

- All overloadable operators are available
- `pack<T> X pack<T>` operations but also `pack<T> X T`
- Type coercion and promotion disabled
`uint8_t(255)+ uint8_t(1)` yields `uint8_t(0)`, not `int(256)`

Comparisons

- `==`, `!=`, `<`, `<=`, `>` and `>=` perform lexical comparisons.
- `eq`, `neq`, `lt`, `gt`, `le` and `ge` as functions return `pack` of boolean.

Other properties

- Models both a `ReadOnlyRandomAccessFusionSequence` and `ReadOnlyRandomAccessRange`
- `at_c<i>(p)` or `p[i]` can be used to access the *i*-th element, but is usually slow (`at_c` is faster)

pack API

Memory access

Memory must be aligned on `sizeof(T)*N` to load/store a `pack<T, N>` from or to a `T*`.
Errors leads to undefined behaviors.

Examples

pack API


Memory access

Memory must be aligned on `sizeof(T)*N` to load/store a `pack<T, N>` from or to a `T*`. Errors leads to undefined behaviors.

Examples

`load< pack<T, N> >(p, i)` loads pack at aligned address `p + i*N`

Main Memory



`load<pack<float>>(0x10,0)`


pack API

Memory access

Memory must be aligned on `sizeof(T)*N` to load/store a `pack<T, N>` from or to a `T*`. Errors leads to undefined behaviors.

Examples

`load< pack<T, N>, Offset>(p, i)` loads pack at address $p + i*N + \text{Offset}$, $p + i$ must be aligned.



pack API

Memory access

Memory must be aligned on `sizeof(T)*N` to load/store a `pack<T, N>` from or to a `T*`.
Errors leads to undefined behaviors.

Examples

`store(p, i, pk)` stores pack `pk` at aligned address `p + i*N`

pack as a proto entity

Rationale

- Most SIMD ISA have fused operations (FMA, etc...)
- We want to write simple code but yet get best performances out of these
- We need lazy evaluation : proto to the rescue !

Advantage

- All expressions, even those involving functions, generate template expressions that are evaluated on assignment or in the conversion operator
- $a * b + c$ is mapped to `fma(a, b, c)`
 $a + b * c$ is mapped to `fma(b, c, a)`
 $!(a < b)$ is mapped to `is_nle(a, b)`
- the optimisation system is open for extensions

Extra arithmetic, bitwise and ieee operations, predicates

Arithmetic

- saturated arithmetic
- float/int conversion
- round, floor, ceil, trunc
- sqrt, hypot
- average
- random
- min/max
- rounded division and remainder

Bitwise

- select
- andnot, ornot
- popcnt
- ffs
- ror, rol
- rshr, rshl
- twopower

IEEE

- ilogb, frexp
- ldexp

- next/prev
- ulpdist

Predicates

- comparison with zero
- negation of comparison
- is_unord, is_nan, is_invalid
- is_odd, is_even
- majority

Reduction and SWAR operations

Reduction

- any, all
- nbtrue
- minimum/maximum, posmin/posmax
- sum
- product, dot product

SWAR

- group/split
- splatted reduction
- cumsum
- sort

native<T, X> : SIMD register of T on arch. X

Semantic

- like `pack` but Plain Old Data and all operations and functions return values and not expression templates.
- `x` characterizes the register type, not the instructions available. Only one tag for all SSE variants.
- It is the interface that must be used to extend the library.

Example

```
native<float, tag::sse_>      wraps a __m128
native<uint8_t, tag::sse_>   wraps a __m128i
native<double, tag::avx_>    wraps a __m256d
native<float, tag::altivec_> wraps a __vector float
```

native<T, X> : SIMD register of T on arch. X

Software fallback

- `tag::none_<N>` is a software-emulated SIMD architecture with a register size of `N` bytes
- It is used as fallback when no satisfying SIMD architecture is found
- Thanks to this, code can degrade well and remain portable.
- Default native type when no SIMD is found : `native<T, tag::none_<8> >`

RGB to grayscale

Scalar version

```
float const *red, *green, *blue;
float* result;

for(std::size_t i = 0; i != height*width; ++i)
    result[i] = 0.3f * red[i] + 0.59f * green[i] + 0.11f * blue[i];
```

SIMD version

```
std::size_t N = meta::cardinal_of<pack<float>>::value;
for(std::size_t i = 0; i != height*width/N; ++i)
{
    pack<float> r = load< pack<float> >(red, i);
    pack<float> g = load< pack<float> >(green, i);
    pack<float> b = load< pack<float> >(blue, i);

    pack<float> res = 0.3f * r + 0.59f * g + 0.11f * b;
    store(res, result, i);
}
```

Easy enough, but what if...

- ... I've got interleaved RGB or RGBA?
- ... I've got 8-bit integers and not floats?

Sounds more complicated, we'll see that later.

Talk Layout

- 1 Introduction
- 2 Interface
- 3 SIMD and STL**
- 4 SIMD Specific Idioms
- 5 Conclusion

Operations vs Data

Where/How to store our data ?

- SIMD operations require data to operate onto
- Usual approaches force a specific container type onto users
- Not generic enough

A better approach

- SIMD compliant allocators
- SIMD Range and Iterators over ContiguousRange
- Adapt our SIMD classes to work with a subset of STD algorithms

Operations vs Data

Where/How to store our data ?

- SIMD operations require data to operate onto
- Usual approaches force a specific container type onto users
- Not generic enough

A better approach

- SIMD compliant allocators
- SIMD Range and Iterators over ContiguousRange
- Adapt our SIMD classes to work with a subset of STD algorithms

SIMD allocators

Rationale

- Allow containers to handle memory in a SIMD compliant way
- Handles alignment of memory
- Handles padding of memory

Example

```
std::vector<float, simd::allocator<float> > v(173);  
  
assert( simd::is_aligned(&v[0]) );
```


From Range to SIMDRange

Iterator interface

- Boost.SIMD provides `simd::begin()/simd::end()`
- Turn iterators into SIMD iterators returning pack
- Take a regular range, iterate over it in SIMD


Example

From Range to SIMDRange

Iterator interface

- Boost.SIMD provides `simd::begin()/simd::end()`
- Turn iterators into SIMD iterators returning pack
- Take a regular range, iterate over it in SIMD

Example



From Range to SIMDRange

Iterator interface

- Boost.SIMD provides `simd::begin()/simd::end()`
- Turn iterators into SIMD iterators returning pack
- Take a regular range, iterate over it in SIMD

Example

```
std::vector<float, simd::allocator<float> > v(1024);  
pack<float> x,z;  
  
x = std::accumulate( simd::begin(v.begin())  
                    , simd::end(v.end())  
                    , z  
                    );
```

From Range to SIMDRange

Iterator interface

- Boost.SIMD provides `simd::begin()/simd::end()`
- Turn iterators into SIMD iterators returning `pack`
- Take a regular range, iterate over it in SIMD

Example

```
std::vector<float, simd::allocator<float> > v(1024);  
pack<float> x,z;  
  
x = boost::accumulate(simd::range(v), z);
```

From Range to SIMDRange

Iterator interface

- native and pack provides `begin()/end()`
- Directly usable in STD algorithms
- Directly usable in Boost.Range algorithms

Example

```
pack<float> x(1,2,3,4);  
  
float k = std::accumulate(x.begin(), x.end(), 0.f);
```

SIMD values as Range

Putting everything together

```
std::vector<float, simd::allocator<float> > v(1024);  
pack<float> x,z;  
float r;  
  
x = boost::accumulate(simd::range(v), z);  
r = std::accumulate(x.begin(), x.end(), 0.f);
```


SIMD values as Range

Putting everything together - Better version

```
std::vector<float, simd::allocator<float> > v(1024);  
float r;  
  
r = sum(accumulate(simd::range(v), pack<float>()));
```

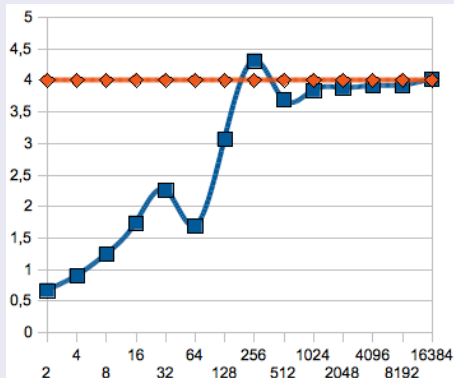
SIMD values as Range

`std::accumulate` speed-up for double



SIMD values as Range

`std::accumulate` speed-up for float



SIMD Range and generic SIMD/scalar code

Back to RGB2Grey

```

template<class RangeIn, class RangeOut> inline void
rgb2grey( RangeOut result, RangeIn red, RangeIn green, RangeIn blue )
{
    typedef typename RangeIn::iterator   in_iterator;
    typedef typename RangeOut::iterator  iterator;
    typedef typename iterator_value<iterator>::type type;

    iterator br = result.begin(), er = result.end();
    in_iterator r = red.begin();
    in_iterator g = green.begin();
    in_iterator b = blue.begin();

    while( br != er )
    {
        type rv = load< type >(r, 0);
        type gv = load< type >(g, 0);
        type bv = load< type >(b, 0);
        type res = 0.3f * rv + 0.59f * gv + 0.11f * bv;
        store(res, br, 0);
        br++; r++; g++; b++;
    }
}

```

What's Missing ?

Integrated SIMD support

- Most STD algorithms should be specialized to be run in one scoop
- Can we have a Boost.Range adaptor like `simd(r)` ?
- Support for shifted Range using `load<T,N>`

Some SIMD mind teasers

- SIMD `find` ?
- SIMD `sort` ?
- Accelerating stuff like `copy` ?

Talk Layout

- 1 Introduction
- 2 Interface
- 3 SIMD and STL
- 4 SIMD Specific Idioms**
- 5 Conclusion

Boolean values in SIMD

The Problem

```
pack<float> x(1,2,3,4);  
pack<float> c(2.5);  
  
cout << lt(x,c) << endl;
```

Boolean values in SIMD

The Problem

```
pack<float> x(1,2,3,4);  
pack<float> c(2.5);  
  
cout << lt(x,c) << endl;  
(( Nan Nan 0 0))
```

Boolean values in SIMD

The Problem

```
pack<float> x(1,2,3,4);  
pack<float> c(2.5);  
  
cout << lt(x,c) << endl;  
(( Nan Nan 0 0))
```

The Solution

`True<T>()` which returns a proper `true` value w/r to `T`
`False<T>()` which returns a proper `false` value w/r to `T`

Conditional in SIMD

Example

```
// Scalar code
if( x > 4 )
    y = 2*x;
else
    z = 1.f/x

// SIMD code
// ???
```

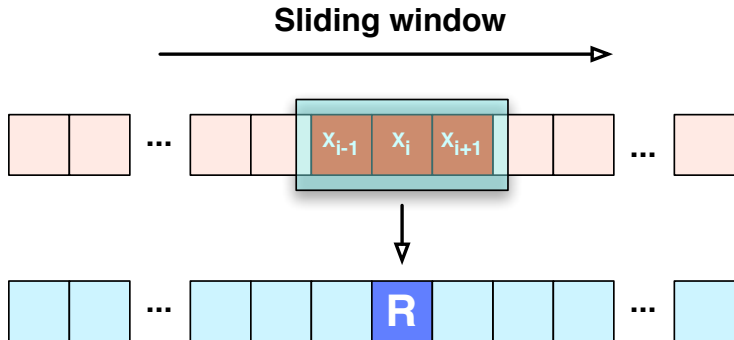

Conditional in SIMD

Example

```
// Scalar code
if( x > 4 )
    y = 2*x;
else
    z = 1.f/x


// SIMD code
y = where( gt(x, 4), 2*x, y);
z = where( gt(x, 4), z, 1.f/x);
```

Motivation



$$R[i] = 1/3 * (x[i-1] + x[i] + x[i+1])$$

Getting there ...




$$R[4*i+0] = 1/3 * (x[4*i-1] + x[4*i+0] + x[4*i+1])$$

$$R[4*i+1] = 1/3 * (x[4*i+0] + x[4*i+1] + x[4*i+2])$$

$$R[4*i+2] = 1/3 * (x[4*i+1] + x[4*i+2] + x[4*i+3])$$

$$R[4*i+3] = 1/3 * (x[4*i+2] + x[4*i+3] + x[4*i+4])$$

The vector solution



$$VR = 1/3 * (\text{load}<-1>(vx) + \text{load}<0>(vx) + \text{load}<1>(vx))$$

The vector solution

SIMD/Scalar version

```
template<class RangeIn, class RangeOut>
inline void average( RangeOut result, RangeIn input )
{
    typedef typename RangeIn::iterator  in_iterator;
    typedef typename RangeOut::iterator iterator;
    typedef typename iterator_value<iterator>::type type;

    iterator br = result.begin(), er = result.end();
    in_iterator data = input.begin();

    br++; er--;
    while( br != er )
    {
        type xm1 = load<type, -1>(data,i);
        type x   = load<type>(data,i);
        type xp1 = load<type, +1>(data,i);

        store(res,i,0) = 1.f/3 * (xm1 + x + xp1);
    }
}
```

Back to RGB2Grey

8-bit RGB

```
static const std::size_t N = meta::cardinal_of< pack<uint8_t> >::value;
for(std::size_t i = 0; i != height*width/N; ++i)
{
    pack<uint8_t> r = load< pack<uint8_t> >(red, i);
    pack<uint8_t> g = load< pack<uint8_t> >(green, i);
    pack<uint8_t> b = load< pack<uint8_t> >(blue, i);

    pack<uint8_t> res = uint8_t(77) * r / uint8_t(255) + uint8_t(150)
        * g / uint8_t(255) + uint8_t(28) * b / uint8_t(255);
    store(res, result, i);
}
```

Back to RGB2Grey

Promote the pack

```
uint16_t r_coeff = 77;  
uint16_t g_coeff = 150;  
uint16_t b_coeff = 28;  
uint16_t div_coeff = 255;
```

```
pack<uint16_t> r1, r2, g1, g2, b1, b2;  
tie(r1, r2) = split(r);  
tie(g1, g2) = split(g);  
tie(b1, b2) = split(b);
```

```
pack<uint16_t> res1 = (r_coeff * r1 + g_coeff * g1 + b_coeff * b1) / div_coeff;  
pack<uint16_t> res2 = (r_coeff * r2 + g_coeff * g2 + b_coeff * b2) / div_coeff;  
pack<uint8_t> res = group(res1, res2);
```

Talk Layout

- 1 Introduction
- 2 Interface
- 3 SIMD and STL
- 4 SIMD Specific Idioms
- 5 Conclusion

Overview of Boost.SIMD

Our goals

- Bring SIMD programming to a usable state
- if we have `boost.atomic`, why not `boost.simd` ?
- Be attractive by being nice with the rest of C++

What we achieved

- Leveraging what we learned in NT^2
- Demonstrated some impacts in term of performance
- Made using SIMD almost as simple than scalar

Upcoming works

Google Summer of Code 2011

- Cleanign up the mess and boostify it
- Improve STL/Boost compatibility
- **Wanted: Applications so we can have real life examples in the library**

Thanks for your attention