

~~An Introduction to~~ Lock-free Programming
Continued

Tony Van Eerd
BoostCon 2011

1.
Summary of
Lockfree BoostCon 2010

2 hours of content that was almost presented in 90min, condensed to 30min
– or less –

“Use Locks!”

A Guide to Threaded Coding

A Guide to Threaded Coding

1. Forget what you learned in Kindergarten

A Guide to Threaded Coding

1. Forget what you learned in Kindergarten
(stop Sharing)

A Guide to Threaded Coding

1. Forget what you learned in Kindergarten
(stop Sharing)
2. Use Locks

A Guide to Threaded Coding

1. Forget what you learned in Kindergarten
(stop Sharing)
2. Use Locks
3. Measure

A Guide to Threaded Coding

1. Forget what you learned in Kindergarten
(stop Sharing)
2. Use Locks
3. Measure
4. Measure

A Guide to Threaded Coding

1. Forget what you learned in Kindergarten
(stop Sharing)
2. Use Locks
3. Measure
4. Measure
5. Change your Algorithm

A Guide to Threaded Coding

1. Forget what you learned in Kindergarten
(stop Sharing)
2. Use Locks
3. Measure
4. Measure
5. Change your Algorithm
6. GOTO 1

A Guide to Threaded Coding

1. Forget what you learned in Kindergarten
(stop Sharing)
2. Use Locks
3. Measure
4. Measure
5. Change your Algorithm
6. GOTO 1

- ∞. Lock-free

A Guide to Threaded Coding

“Lock-free coding is the last thing you want to do.”

1. Forget what you learned in Kindergarten
(stop Sharing)
2. Use Locks
3. Measure
4. Measure
5. Change your Algorithm
6. GOTO 1

- ∞. Lock-free

Example 1

Bad

Thread 1,2,3...:

```
static bool locked;  
  
if ( !locked )  
{  
    locked = true;  
    do_exclusive_stuff();  
}  
else...
```

Example 1

Bad

Thread 1,2,3...:

```
static bool locked;  
  
if ( !locked )  
{  
  < ? >  
  locked = true;  
  do_exclusive_stuff();  
}  
else...
```

“Read between the lines”

Example 1 - Atomically

Bad

Thread 1,2,3...:

```
static bool locked;
```

```
if ( !locked )  
{  
    locked = true;  
    do_exclusive_stuff();  
}  
else...
```

Atomically

How?

Example 1 - Atomically

Good - “CAS”

Thread 1,2,3...:

```
static std::atomic<bool> locked;  
  
if (locked.compare_exchange(false, true))  
{  
    do_exclusive_stuff();  
}  
else...
```

(CAS = Compare-And-Set == Compare-And-Swap == Test-And-Set == compare_exchange)

Example 1 - Atomically

Bad

```
static bool locked;

if ( !locked )
{
    locked = true;
    do_exclusive_stuff();
}
else...
```

Good - CAS

```
static std::atomic<bool> locked;

if (locked.compare_exchange(false, true))
{
    do_exclusive_stuff();
}
else...
```

Example 1 - Atomically

Bad

```
static bool locked;

if ( !locked )
{
    locked = true;
    do_exclusive_stuff();
}
else...
```

Good - CAS

```
static std::atomic<bool> locked;

if (locked.compare_exchange(false, true))
{
    do_exclusive_stuff();
}
else...
```

Simple, Right?...

Example 2

Bad

Thread 1:

```
data.x = .....;  
data.y = .....;  
data_ready = true;
```

Thread 2:

```
if (data_ready)  
{  
    a = data.x;  
    b = data.y;  
    ...  
}
```

Example 2

Thread 1:

Your Code:

```
data.x = .....;  
data.y = .....;  
data_ready = true;
```

Example 2

Thread 1:

Your Code:

```
data.x = .....;
data.y = .....;
data_ready = true;
```



Evil (or Smart?) CPU(s):

?



```
data_ready = true;
data.x = .....;
data.y = .....;
```

That was Then:

speed of CPU == speed of RAM

That was Then:

speed of CPU == speed of RAM

This is Now:

speed of CPU >>> speed of RAM
100x!

CPU and compiler optimization/reordering rule:

Do any optimizations you'd like, as long as it doesn't change how the program works.*

*....assuming a single-threaded program.

Example 2

Thread 1:

Your Code:

```
data.x = .....;
data.y = .....;
data_ready = true;
```



Evil/Smart/Fast CPU(s):

?



```
data_ready = true;
data.x = .....;
data.y = .....;
```

Example 2

Thread 1:

Your Code:

```
data.x = .....;  
data.y = .....;  
data_ready = true;
```



Evil/Smart/Fast CPU(s):

Shift
Happens



```
data_ready = true;  
data.x = .....;  
data.y = .....;
```

Example 2

Thread 1:

Your Code:



Evil/Smart/Fast CPU(s):

```
data.x = .....;
data.y = .....;
data_ready = true;
```

Shift
Happens



```
data_ready = true;
data.x = .....;
data.y = .....;
```

Thread 2:

Your Code:



Evil/Smart/Fast CPU(s):

```
if (data_ready)
{
    a = data.x;
    b = data.y;
    ...
}
```

Can't reorder an if!
(?)

Example 2

Thread 1:

Your Code:

```
data.x = .....;
data.y = .....;
data_ready = true;
```



Evil/Smart/Fast CPU(s):

Shift
Happens



```
data_ready = true;
data.x = .....;
data.y = .....;
```

Thread 2:

Your Code:

```
if (data_ready)
{
    a = data.x;
    b = data.y;
    ...
}
```



Evil/Smart/Fast CPU(s):

Shift
Happens



```
tmp = read(data);
if (data_ready)
{
    a = tmp.x;
    b = tmp.y;
}
```

Example 2

Bad

```
static bool data_ready;
```

Thread 1:

```
data.x = .....;  
data.y = .....;  
data_ready = true;
```

Thread 2:

```
if (data_ready)  
{  
    a = data.x;  
    b = data.y;  
    ...  
}
```

Example 2

Good - C++0x

```
static std::atomic<bool> data_ready;
```

Thread 1:

```
data.x = .....;  
data.y = .....;  
data_ready = true;
```

Thread 2:

```
if (data_ready)  
{  
    a = data.x;  
    b = data.y;  
    ...  
}
```

Example 2

Good - C++0x

```
static std::atomic<bool> data_ready;
```

Thread 1:

```
data.x = .....;  
data.y = .....;  
data_ready = true;
```

Thread 2:

```
if (data_ready)  
{  
    a = data.x;  
    b = data.y;  
    ...  
}
```


Example 2

Good - C++0x

```
static std::atomic<bool> data_ready;
```

Thread 1:

```
data.x = .....;
data.y = .....;
data_ready = true;
q = 10;
```

Thread 2:

```
w = r;
if (data_ready)
{
    a = data.x;
    b = data.y;
    ...
}
```

Example 2

Good - C++0x

```
static std::atomic<bool> data_ready;
```

Thread 1:

```
data.x = .....;  
data.y = .....;  
data_ready = true;
```

Thread 2:

```
if (data_ready)  
{  
    a = data.x;  
    b = data.y;  
    ...  
}
```

Example 2

Good - C++0x

```
static std::atomic<bool> data_ready;
```

Thread 1:

```
data.x = .....;  
data.y = .....;  
data_ready.store(true, std::memory_order_release);
```

Thread 2:

```
if (data_ready)  
{  
    a = data.x;  
    b = data.y;  
    ...  
}
```

Example 2

Good - C++0x

```
static std::atomic<bool> data_ready;
```

Thread 1:

```
data.x = .....;
data.y = .....;
data_ready.store(true, std::memory_order_release);
```

release: before means before

Thread 2:

```
if (data_ready)
{
    a = data.x;
    b = data.y;
    ...
}
```

Example 2

Good - C++0x

```
static std::atomic<bool> data_ready;
```

Thread 1:

```
data.x = .....;
data.y = .....;
data_ready.store(true, std::before_means_before);
```

release: before means before

Thread 2:

```
if (data_ready)
{
    a = data.x;
    b = data.y;
    ...
}
```

Example 2

Good - C++0x

```
static std::atomic<bool> data_ready;
```

Thread 1:

```
data.x = .....;  
data.y = .....;  
data_ready.store(true, std::memory_order_release);
```

Thread 2:

```
if (data_ready)  
{  
    a = data.x;  
    b = data.y;  
    ...  
}
```

Example 2

Good - C++0x

```
static std::atomic<bool> data_ready;
```

Thread 1:

```
data.x = .....;
data.y = .....;
data_ready.store(true, std::memory_order_release);
```

Thread 2:

```
if (data_ready)
{
    a = data.x;
    b = data.y;
    ...
}
```

```
tmp = read(data);
if (data_ready)
{
    a = tmp.x;
    b = tmp.y;
    ...
}
```

Example 2

Good - C++0x

```
static std::atomic<bool> data_ready;
```

Thread 1:

```
data.x = .....;
data.y = .....;
data_ready.store(true, std::memory_order_release);
```

Thread 2:

```
if (data_ready.load(std::memory_order_acquire))
{
    a = data.x;
    b = data.y;
    ...
}
```


Example 2

Good - C++0x

```
static std::atomic<bool> data_ready;
```

Thread 1:

```
data.x = .....;
data.y = .....;
data_ready.store(true, std::memory_order_release);
```

Thread 2:

```
if (data_ready.load(std::memory_order_acquire))
{
    a = data.x;
    b = data.y;
    ...
}
```

acquire: after means after

Example 2

Good - C++0x

```
static std::atomic<bool> data_ready;
```

Thread 1:

```
data.x = ....;  
data.y = ....;  
data_ready.store(true, std::memory_order_release);
```

release: before means before

Thread 2:

```
if (data_ready.load(std::memory_order_acquire))  
{  
    a = data.x;  
    b = data.y;  
    ...  
}
```

acquire: after means after

Example 2

Good - C++0x

```
static std::atomic<bool> data_ready;
```

Thread 1:

CPU 1

```
data.x = ....;
```

```
data.y = ....;
```

```
data_ready.store(true, std::memory_order_release);
```

release: before means before

Thread 2:

```
if (data_ready.load(std::memory_order_acquire))
```

```
{
```

```
    a = data.x;
```

```
    b = data.y;
```

```
    ...
```

```
}
```

CPU 2

acquire: after means after

Example 2

Good - C++0x

```
static std::atomic<bool> data_ready;
```

Thread 1:

CPU 1

release: *before means before*

```
A data.x = ....;
  data.y = ....;
  data_ready.store(true, std::memory_order_release);
```

Thread 2:

```
B if (data_ready.load(std::memory_order_acquire))
  {
    a = data.x;
    b = data.y;
    ...
  }
```

CPU 2

acquire: *after means after*

Example 2

Good - C++0x

```
static std::atomic<bool> data_ready;
```

Thread 1:

```
A data.x = .....;
    data.y = .....;
    data_ready = true;
```

Thread 2:

```
B if (data_ready)
    {
        a = data.x;
        b = data.y;
        ...
    }
```

Example 2

Good - C++0x

```
static std::atomic<bool> data_ready;
```

Thread 1:

A	<pre>data.x =; data.y =; data_ready = true;</pre>	<pre>data.x =; data.y =; data_ready.store(true, release);</pre>
----------	---	---

Thread 2:

B	<pre>if (data_ready) { a = data.x; b = data.y; ... }</pre>	<pre>if(data_ready.load(acquire)) { a = data.x; b = data.y; ... }</pre>
----------	--	---

Example 2 – *Happens Before*

Good – **C++0x**



```
static std::atomic<bool> data_ready;
```

Thread 1:

```
A data.x = ....;
    data.y = ....;
    data_ready = true;
```

```
data.x = ....;
data.y = ....;
data_ready.store(true, release);
```

Thread 2:

```
B if (data_ready)
    {
        a = data.x;
        b = data.y;
        ...
    }
```

```
if(data_ready.load(acquire))
{
    a = data.x;
    b = data.y;
    ...
}
```

Happens-before

An operation A *happens-before* an operation B if:

A is performed on the same thread as B, and A is before B in program order,

or

A synchronizes-with B,

or

A happens-before some other operation C, and C happens-before B.


```
data.x = .....; // A  
data.y = .....; // B           // A  
data_ready.store(true, release); // B
```

Happens-before

An operation A *happens-before* an operation B if:

A is performed on the same thread as B, and A is before B in program order,

or

A synchronizes-with B,

or

A happens-before some other operation C, and C happens-before B.

```
data.x = .....; // A // A
data.y = .....; // B // C
data_ready.store(true, release); // B // B
```

Happens-before

An operation A *happens-before* an operation B if:

A is performed on the same thread as B, and A is before B in program order,

or

A synchronizes-with B,

or

A happens-before some other operation C, and C happens-before B.

```
data.x = .....;
data.y = .....;
data_ready.store(true, release);
```

Happens-before

An operation A *happens-before* an operation B if:

A is performed on the same thread as B, and A is before B in program order,

or

A synchronizes-with B,

or

A happens-before some other operation C, and C happens-before B.

```
data.x = ....;
data.y = ....;
data_ready.store(true, release);
```

Happens-before

An operation *A* *happens-before* an operation *B* if:

A is performed on the same thread as *B*, and *A* is before *B* in program order,

or

A synchronizes-with B,

or

A happens-before some other operation *C*, and *C* happens-before *B*.

Synchronizes-with

An operation *A* *synchronizes-with* an operation *B* if

A is a store to some atomic variable *m*, with an ordering of `std::memory_order_release`, or `std::memory_order_seq_cst`,

and

B is a load from the same variable *m*, with an ordering of `std::memory_order_acquire` or `std::memory_order_seq_cst`,

and

B reads the value stored by *A*.

```
data.x = ....;  
data.y = ....;  
data_ready.store(true, std::memory_order_release); //A
```

Happens-before

An operation A *happens-before* an operation B if:

A is performed on the same thread as B, and A is before B in program order,

or

A synchronizes-with B,

or

A happens-before some other operation C, and C happens-before B.

Synchronizes-with

An operation A *synchronizes-with* an operation B if

A is a **store** to some **atomic** variable m, with an ordering of std::**memory_order_release**, or std::memory_order_seq_cst,

and

B is a load from the same variable m, with an ordering of std::memory_order_acquire or std::memory_order_seq_cst,

and

B reads the value stored by A.

```
data.x = ....;
data.y = ....;
data_ready.store(1, release);
```

Happens-before

An operation A *happens-before* an operation B if:

A is performed on the same thread as B, and A is before B in program order,

or

A synchronizes-with B,

or

A happens-before some other operation C, and C happens-before B.

Synchronizes-with

An operation A *synchronizes-with* an operation B if

A is a **store** to some **atomic** variable m, with an ordering of `std::memory_order_release`, or `std::memory_order_seq_cst`,

and

B is a load from the same variable m, with an ordering of `std::memory_order_acquire` or `std::memory_order_seq_cst`,

and

B reads the value stored by A.

<pre>data.x =; data.y =; data_ready.store(1, release);</pre>	<pre>if (data_ready.load(acquire)) { a = data.x; b = data.y; }</pre>
--	--

Happens-before

An operation A *happens-before* an operation B if:

A is performed on the same thread as B, and A is before B in program order,

or

A synchronizes-with B,

or

A happens-before some other operation C, and C happens-before B.

Synchronizes-with

An operation A *synchronizes-with* an operation B if

A is a **store** to some **atomic** variable m, with an ordering of `std::memory_order_release`, or `std::memory_order_seq_cst`,

and

B is a **load** from the same variable m, with an ordering of `std::memory_order_acquire` or `std::memory_order_seq_cst`,

and

B reads the value stored by A.

```
data.x = ....;
data.y = ....;
data_ready.store(1, release);

if (data_ready.load(acquire))
{
    a = data.x;
    b = data.y;
}
```

Happens-before

An operation A *happens-before* an operation B if:

A is performed on the same thread as B, and A is before B in program order,

or

A synchronizes-with B,

or

A happens-before some other operation C, and C happens-before B.

Synchronizes-with

An operation A *synchronizes-with* an operation B if

A is a **store** to some **atomic** variable m, with an ordering of `std::memory_order_release`, or `std::memory_order_seq_cst`,

and

B is a **load from the same variable** m, with an ordering of `std::memory_order_acquire` or `std::memory_order_seq_cst`,

and

B reads the value stored by A.

<pre> data.x =; data.y =; data_ready.store(1, release); </pre>	<pre> if(data_ready.load(acquire)) { a = data.x; b = data.y; } </pre>
---	--

Happens-before

An operation A *happens-before* an operation B if:

A is performed on the same thread as B, and A is before B in program order,

or

A synchronizes-with B,

or

A happens-before some other operation C, and C happens-before B.

Synchronizes-with

An operation A *synchronizes-with* an operation B if

A is a **store** to some **atomic** variable m, with an ordering of `std::memory_order_release`, or `std::memory_order_seq_cst`,

and

B is a **load from the same variable** m, with an ordering of `std::memory_order_acquire` or `std::memory_order_seq_cst`,

and

B reads the value stored by A.

<pre>data.x =; data.y =; data_ready.store(1, release);</pre>	<pre>if (data_ready.load(acquire)) { a = data.x; b = data.y; }</pre>
---	---

Happens-before

An operation A *happens-before* an operation B if:

A is performed on the same thread as B, and A is before B in program order,

or

A synchronizes-with B,

or

A happens-before some other operation C, and C happens-before B.

Synchronizes-with

An operation A *synchronizes-with* an operation B if

A is a **store** to some **atomic** variable m, with an ordering of `std::memory_order_release`, or `std::memory_order_seq_cst`,

and

B is a **load from the same variable** m, with an ordering of `std::memory_order_acquire` or `std::memory_order_seq_cst`,

and

B reads the value stored by A.

<pre> data.x =; data.y =; data_ready.store(1, release); </pre>	<pre> if (data_ready.load(acquire)) { a = data.x; b = data.y; } </pre>
---	---

Happens-before

An operation A *happens-before* an operation B if:

A is performed on the same thread as B, and A is before B in program order,

or

A synchronizes-with B,

or

A happens-before some other operation C, and C happens-before B.

Synchronizes-with

An operation A *synchronizes-with* an operation B if

A is a **store** to some **atomic** variable m, with an ordering of `std::memory_order_release`, or `std::memory_order_seq_cst`,

and

B is a **load from the same variable** m, with an ordering of `std::memory_order_acquire` or `std::memory_order_seq_cst`,

and

B reads the value stored by A.

A`data.x =;``data.y =;``data_ready.store(1, release);``if (data_ready.load(acquire))``{`**B**`a = data.x;``b = data.y;``}`

Happens-before

An operation **A** *happens-before* an operation **B** if:

A is performed on the same thread as B, and A is before B in program order,

or

A synchronizes-with B,

or

A happens-before some other operation C, and C happens-before B.

Synchronizes-with

An operation A *synchronizes-with* an operation B if

A is a store to some atomic variable m, with an ordering of `std::memory_order_release`, or `std::memory_order_seq_cst`,

and

B is a load from the same variable m, with an ordering of `std::memory_order_acquire` or `std::memory_order_seq_cst`,

and

B reads the value stored by A.

Summary

Atomic (no reading/acting between the lines)

<pre>if (!locked) { locked = true; do_exclusive(); }</pre>	<pre>if (locked.compare_exchange(false, true)) { do_exclusive(); }</pre>
--	--

Happens-before (ensure order when needed)

<pre>data = ...; ready = true; if (ready) { use(data); }</pre>	<pre>data = ...; ready.store(true, memory_order_release); if (ready.load(memory_order_acquire)) { use(data); }</pre>
--	--

Atomic (no reading/acting between the lines)

<pre>if (!locked) { locked = true; do_exclusive(); }</pre>	<pre>if (locked.compare_exchange(false, true)) { do_exclusive(); }</pre>
--	--

Atomic (no reading/acting between the lines)

```
if (locked.compare_exchange(false, true))  
{  
    do_exclusive();  
}
```

Atomic (no reading/acting between the lines)

```
if (locked.compare_exchange(false, true))  
{  
    do_exclusive(data);  
}
```


Atomic + Happens-before

```
if (locked.compare_exchange(false, true, memory_order_acquire))  
{  
    do_exclusive(data);  
}
```

Atomic + Happens-before

```
if (locked.compare_exchange(false, true, memory_order_acquire))  
{  
    do_exclusive(data);  
    locked.store(false, memory_order_release);  
}
```

Atomic + Happens-before

```
std::atomic<bool> locked;
```

```
if (locked.compare_exchange(false, true, memory_order_acquire))  
{  
    do_exclusive(data);  
    locked.store(false, memory_order_release);  
}
```

Atomic + Happens-before

```
std::atomic<bool> locked;
```

```
if (locked.compare_exchange(false, true, memory_order_acquire))  
{  
    do_exclusive(data);  
    locked.store(false, memory_order_release);  
}
```

No reordering

No “Reading between the lines”

Reading between the Tokens

Reading between the Tokens

```
count++;
```

Reading between the Tokens

```
count++;
```

=>

```
count = count + 1;
```

Reading between the Tokens

```
count++;
```

=>

```
register reg = read_memory(&count);  
reg = reg + 1;  
write_memory(&count, reg);
```


Reading between the Tokens

```
count++;
```

=>

```
register reg = read_memory(&count);  
reg = reg + 1;  
write_memory(&count, reg);
```

“Read between the lines”

Reading between the Tokens

```
count++;
```

“Read between the... tokens”

=>

```
register reg = read_memory(&count);  
reg = reg + 1;  
write_memory(&count, reg);
```

“Read between the lines”

Reading between the Tokens

```
count++;
```

“Read between the... tokens”

=>

```
register reg = read_memory(&count);  
reg = reg + 1;  
write_memory(&count, reg);
```

“Read between the lines”

Solution?...

Atomic Increment

```
std::atomic<int> count;
```

```
count++;
```

Atomic Increment

```
std::atomic<int> count;
```

```
count++;
```

=>

????

```
  /\_/\
 (  @.@" )
  =^=
```

Atomic Increment

```
std::atomic<int> count;
```

```
count++;
```

=>

```
// lock or TM magic?
```

```
atomically {  
    register reg = read_memory(&count);  
    reg = reg + 1;  
    write_memory(&count, reg);  
}
```

Atomic Increment

```
std::atomic<int> count;
```

```
count++;
```

=>

```
// processor magic?  
__asm {  
    LOCK inc @count  
}
```

Atomic Increment

```
std::atomic<int> count;
```

```
count++;
```

=>

```
//?
```

```
register reg = read_memory(&count);
```

```
reg = reg + 1;
```

```
write_memory(&count, reg);
```



Thread2: count++;

Atomic Increment

```
std::atomic<int> count;
```

```
count++;
```

=>

```
//?
```

```
register old = read_memory(&count);
```

```
register new = old + 1;
```

```
count.compare_exchange(old, new); //?
```

Atomic Increment

```
std::atomic<int> count;
```

```
count++;
```

=>

```
retry:
```

```
register old = read_memory(&count);
```

```
register new = old + 1;
```

```
if ( ! count.compare_exchange(old, new) )
```

```
    goto retry;
```

Atomic Increment

```
std::atomic<int> count;
```

```
count++;
```

=>

```
do {  
    register old = read_memory(&count);  
    register new = old + 1;  
} while ( ! count.compare_exchange(old, new) );
```

Atomic Increment

```
std::atomic<int> count;
```

```
count++;
```

=>

```
do {  
    int old = count.load(std::memory_order_relaxed);  
    int newc = old + 1;  
} while ( ! count.compare_exchange(old, newc));
```

CAS Loop

```
do {  
    int old = count.load(std::memory_order_relaxed);  
    int newc = old + 1;  
} while ( ! count.compare_exchange(old, newc));
```

CAS Loop

Read global

```
do {  
    int old = count.load(std::memory_order_relaxed);  
    int newc = old + 1;  
} while ( ! count.compare_exchange(old, newc) );
```

CAS Loop

```
do {  
  Act local int old = count.load(std::memory_order_relaxed);  
             int newc = old + 1;  
} while ( ! count.compare_exchange(old, newc));
```

CAS Loop

```
do {  
    int old = count.load(std::memory_order_relaxed);  
    int newc = old + 1;  
} while ( ! count.compare_exchange(old, newc) );
```

CAS global

CAS Loop

```
do {  
    int old = count.load(std::memory_order_relaxed);  
    int newc = old + 1;  
} while ( ! count.compare_exchange(old, newc));
```

***“Think Global
Act Local”***

CAS Loop

```
do {  
    int old = count.load(std::memory_order_relaxed);  
    int newc = old + 1;  
} while ( ! count.compare_exchange(old, newc) );
```

*“Act Local
CAS Global”*

“CAS”

C++0x

- std::memory_order_seq_cst
- std::memory_order_acq_rel
- std::memory_order_acquire
- std::memory_order_release
- std::memory_order_consume
- std::memory_order_relaxed


compare_exchange_strong(volatile type* x, type* expected, type desired, success_order, failure_order)
compare_exchange_weak(volatile type* x, type* expected, type desired, success_order, failure_order)

CAS – with barriers/memory_order assumed (store/load/release/acquire/whatever)

2.



Data Structures

“Easy” - Push a Value onto a Lock-free Stack



```
void push(Val val)
{
}
```

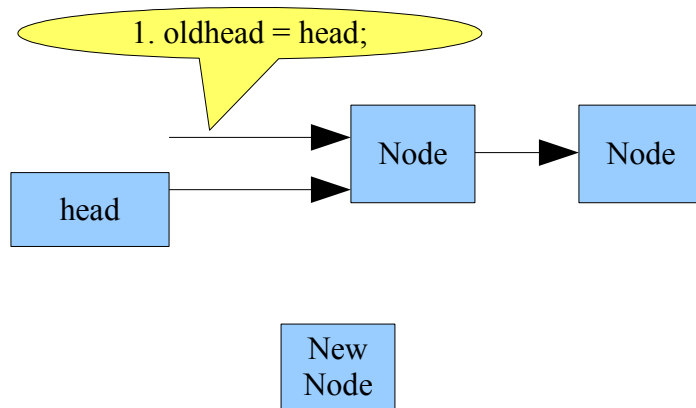
“Easy” - Push a Value onto a Lock-free Stack



```
0. newhead = new Node;
```

```
void push(Val val)
{
    Node * newhead = new Node(val);    // 0
}
```


“Easy” - Push a Value onto a Lock-free Stack



0. newhead = new Node;


```
void push(Val val)
{
    Node * newhead = new Node(val);    // 0
    Node * oldhead = stack.head;       // 1 Read
}
```


“Easy” - Push a Value onto a Lock-free Stack



```
void push(Val val)
{
    Node * newhead = new Node(val);    // 0
    Node * oldhead = stack.head;      // 1 Read
    newhead->next = oldhead;          // 2 Act Locally
}
```


“Easy” - Push a Value onto a Lock-free Stack



```
void push(Val val)
{
    Node * newhead = new Node(val); // 0
    Node * oldhead = stack.head;    // 1 Read
    newhead->next = oldhead;        // 2 Act Locally
    stack.head = newhead;          // 3 Write
}
```

*Write, right?
Wrong.*

“Easy” - Push a Value onto a Lock-free Stack




```
void push(Val val)
{
    Node * newhead = new Node(val); // 0
    do
    {
        Node * oldhead = stack.head; // 1 Read
        newhead->next = oldhead;      // 2 Act Locally
    }                                 // 3 CAS Globally:
    while(!stack.head.CAS(oldhead, newhead));
}
```

Take a good look – that's as easy as it gets!

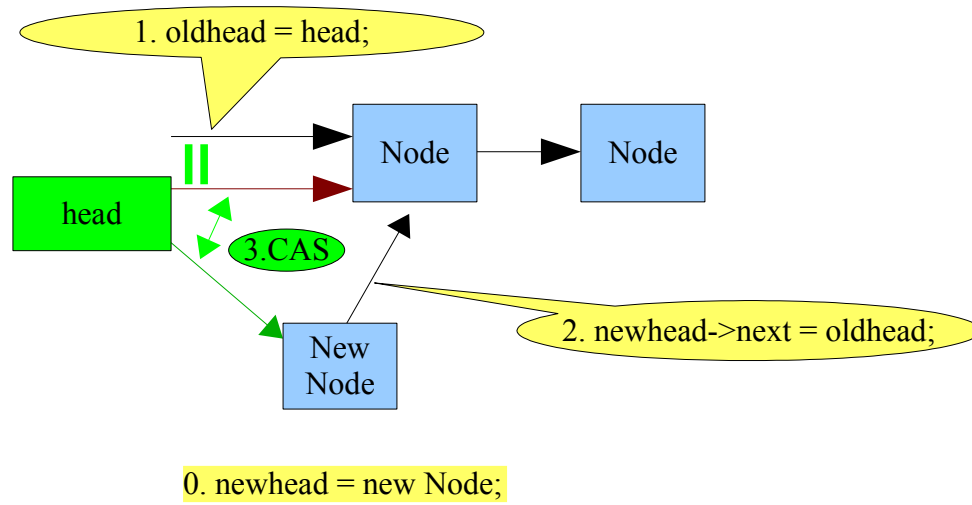
P.S. What memory ordering is required, and where? (and why?)

Act Local




0. `newhead = new Node;`

CAS Global




“Hard” - Pop a Value from a Lock-free Stack




```
Val pop()  
{  
}
```

“Hard” - Pop a Value from a Lock-free Stack




```
Val pop()  
{  
    Node * oldhead = stack.head;    // 1 Read  
}
```

“Hard” - Pop a Value from a Lock-free Stack



```
Val pop()  
{  
  Node * oldhead = stack.head;    // 1 Read  
  Node * newhead = oldhead->next; // 2 Act Locally  
}
```



“Hard” - Pop a Value from a Lock-free Stack



```
Val pop()  
{  
  Node * oldhead = stack.head;    // 1 Read  
  if (!oldhead)  
    throw StackEmpty(); // or return null, etc  
  Node * newhead = oldhead->next; // 2 Act Locally  
}
```

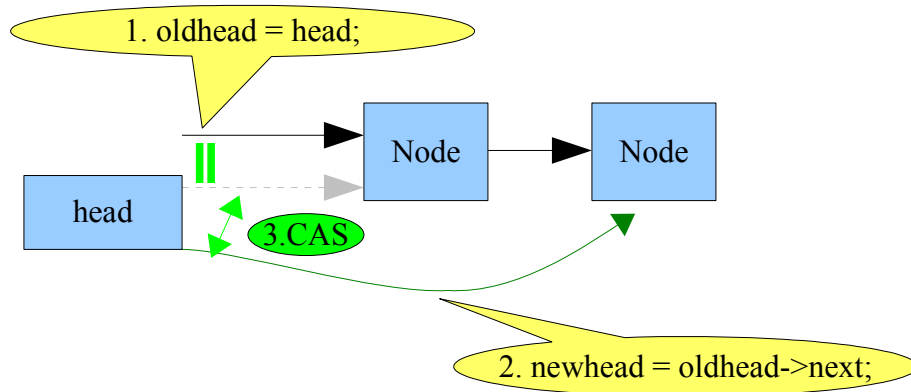


“Hard” - Pop a Value from a Lock-free Stack



```
Val pop()  
{  
    Node * oldhead = stack.head;    // 1 Read  
    if (!oldhead)  
        throw StackEmpty(); // or return null, etc  
    Node * newhead = oldhead->next; // 2 Act Locally  
}
```

“Hard” - Pop a Value from a Lock-free Stack




```
Val pop()  
{  
  Node * newhead;  
  do  
  {  
    Node * oldhead = stack.head;    // 1 Read  
    if (!oldhead)  
      throw StackEmpty(); // or return null, etc  
    newhead = oldhead->next;        // 2 Locally  
  }  
  while (!head.CAS(oldhead, newhead)); // 3 CAS Globally  
  
  Val val = oldhead->val;  
  recycle(oldhead);  
  return val;  
}
```

What's wrong with this code?

In a word:

ABA


A – Pop at Step 2. Prepped for CAS:




```
Val pop()
{
  Node * newhead;
  do
  {
    Node * oldhead = stack.head;
    if (!oldhead)
      throw StackEmpty();
    newhead = oldhead->next;
  }
  while (!head.CAS(oldhead, newhead));

  Val val = oldhead->val;
  delete oldhead;
  return val;
}
```


AB – Now, another thread interrupts, leaves us with:



```
Val pop()  
{  
    Node * newhead;  
    do  
    {  
        Node * oldhead = stack.head;  
        if (!oldhead)  
            throw StackEmpty();  
        newhead = oldhead->next;  
    }  
    while (!head.CAS(oldhead, newhead));  
  
    Val val = oldhead->val;  
    delete oldhead;  
    return val;  
}
```




AB – CAS will fail. Yeah!



```
Val pop()  
{  
  Node * newhead;  
  do  
  {  
    Node * oldhead = stack.head;  
    if (!oldhead)  
      throw StackEmpty();  
    newhead = oldhead->next;  
  }  
  while (!head.CAS(oldhead, newhead));  
  
  Val val = oldhead->val;  
  delete oldhead;  
  return val;  
}
```




A – Pop at Step 2. Prepped for CAS:



```
Val pop()  
{  
  Node * newhead;  
  do  
  {  
    Node * oldhead = stack.head;  
    if (!oldhead)  
      throw StackEmpty();  
    newhead = oldhead->next;  
  }  
  while (!head.CAS(oldhead, newhead));  
  
  Val val = oldhead->val;  
  delete oldhead;  
  return val;  
}
```


AB – Maybe the first Node was deleted...



```
Val pop()
{
    Node * newhead;
    do
    {
        Node * oldhead = stack.head;
        if (!oldhead)
            throw StackEmpty();
        newhead = oldhead->next;
    }
    while (!head.CAS(oldhead, newhead));

    Val val = oldhead->val;
    delete oldhead;
    return val;
}
```


ABA – ...And then reused (recycled by Allocator) and Pushed:



```
Val pop()  
{  
    Node * newhead;  
    do  
    {  
        Node * oldhead = stack.head;  
        if (!oldhead)  
            throw StackEmpty();  
        newhead = oldhead->next;  
    }  
    while (!head.CAS(oldhead, newhead));  
  
    Val val = oldhead->val;  
    delete oldhead;  
    return val;  
}
```



ABA – CAS Succeeds! (Yeah!?):



```
Val pop()  
{  
  Node * newhead;  
  do  
  {  
    Node * oldhead = stack.head;  
    if (!oldhead)  
      throw StackEmpty();  
    newhead = oldhead->next;  
  }  
  while (!head.CAS(oldhead, newhead));  
  
  Val val = oldhead->val;  
  delete oldhead;  
  return val;  
}
```

Our friend CAS has issues!

DWCAS to the Rescue!

```
DWCAS ([__word__ | __word__], oldwide, newwide);
```

What should we do with that?

DWCAS to the Rescue!


[__counter__ | __pointer__]

DWCAS – Pop from a lockfree stack


```
Val pop()
{
  NodePtr newhead;
  do
  {
    NodePtr oldhead = stack.head;
    if (!oldhead)
      throw StackEmpty();
    newhead = oldhead->next;
  }
  while (!DWCAS(&head, oldhead, newhead));

  Val val = oldhead->val;
  delete oldhead;
  return val;
}


NodePtr NodeAlloc(Val val)
{
  NodePtr ptr;
  ptr.ptr = new Node(val);
  ptr.count = atomic_inc(NodePtr::global_count);
  return NodePtr;
}
```



Circular Queue




Single Threaded Circular Queue



```
bool push(int item)
{
    if (tail == head - 1) // full?
        return false;
    array[tail] = item;
    tail = (tail + 1) % size;
    return true;
}

bool pop(int & item)
{
    if (head == tail) // empty?
        return false;
    item = array[head];
    head = (head + 1) % size;
    return true;
}
```


Single Threaded Circular Queue



```
bool push(int item)
{
    if ((tail + 1) % size == head) // full?
        return false;
    array[tail] = item;
    tail = (tail + 1) % size;
    return true;
}

bool pop(int & item)
{
    if (head == tail) // empty?
        return false;
    item = array[head];
    head = (head + 1) % size;
    return true;
}
```


“Atomic”? Circular Queue



```
std::atomic<size_t> head;  
std::atomic<size_t> tail;
```

```
bool push(int item) {  
    if ((tail + 1) % size == head)  
        return false;  
    array[tail] = item;  
    tail = (tail + 1) % size;  
    return true;  
}  
bool pop(int & item) {  
    if (head == tail)  
        return false;  
    item = array[head];  
    head = (head + 1) % size;  
    return true;  
}
```


“Atomic”? Circular Queue



```
std::atomic<size_t> head;  
std::atomic<size_t> tail;
```


```
bool push(int item) {  
    if ((tail + 1) % size == head)  
        return false;  
    array[tail] = item;  
    tail = (tail + 1) % size;  
    return true;  
}  
bool pop(int & item) {  
    if (head == tail)  
        return false;  
    item = array[head];  
    head = (head + 1) % size;  
    return true;  
}
```

Single Threaded Circular Queue... of non-zero ints




```
bool push(int item) {
    if (array[tail] != 0) //occupied => full
        return false;
    array[tail] = item;
    tail = (tail + 1) % size;
    return true;
}
bool pop(int & item) {
    if (array[head] == 0) // empty?
        return false;
    item = array[head];
    array[head] = 0;
    head = (head + 1) % size;
    return true;
}
```

Single Threaded Circular Queue... of non-zero ints



```
bool push(int item) {
    if (array[tail] != 0) //occupied => full
        return false;
    array[tail] = item;
    tail = (tail + 1) % size;
    return true;
}
bool pop(int & item) {
    item = array[head];
    if (item == 0) // empty?
        return false;
    array[head] = 0;
    head = (head + 1) % size;
    return true;
}
```


“Atomic”? Circular Queue... of non-zero ints



```
std::atomic<int> array[size];
```

```
bool push(int item) {  
    if (array[tail].load(memory_order_???) != 0) )  
        return false;  
    array[tail].store(item, memory_order_???) ;  
    tail = (tail + 1) % size;  
    return true;  
}  
bool pop(int & item) {  
    item = array[head].load(memory_order_???) ;  
    if (item == 0)  
        return false;  
    array[head].store(0, memory_order_???) ;  
    head = (head + 1) % size;  
    return true;  
}
```


“Atomic”? Circular Queue... of non-zero ints



```
std::atomic<int> array[size];
```

```
bool push(int item) {  
    if (array[tail].load(memory_order_relaxed) != 0)  
        return false;  
    array[tail].store(item, memory_order_relaxed);  
    tail = (tail + 1) % size;  
    return true;  
}  
bool pop(int & item) {  
    item = array[head].load(memory_order_relaxed);  
    if (item == 0)  
        return false;  
    array[head].store(0, memory_order_relaxed);  
    head = (head + 1) % size;  
    return true;  
}
```



“Atomic”? Circular Queue... of non-zero ints ←



```
std::atomic<int> array[size];
```

```
bool push(int item) {  
    if (array[tail].load(memory_order_relaxed) != 0) {  
        return false;  
    }  
    array[tail].store(item, memory_order_relaxed);  
    tail = (tail + 1) % size;  
    return true;  
}  
bool pop(int & item) {  
    item = array[head].load(memory_order_relaxed);  
    if (item == 0) {  
        return false;  
    }  
    array[head].store(0, memory_order_relaxed);  
    head = (head + 1) % size;  
    return true;  
}
```


SPSC Circular Queue of non-zero ints



```
std::atomic<int> array[size];
```

```
bool push(int item) {  
    if (array[tail].load(memory_order_relaxed) != 0) {  
        return false;  
    }  
    array[tail].store(item, memory_order_relaxed);  
    tail = (tail + 1) % size;  
    return true;  
}  
bool pop(int & item) {  
    item = array[head].load(memory_order_relaxed);  
    if (item == 0) {  
        return false;  
    }  
    array[head].store(0, memory_order_relaxed);  
    head = (head + 1) % size;  
    return true;  
}
```


*SPSC Circular Queue of **non-null pointers***



```
std::atomic<Item *> array[size];
```

```
bool push(Item * itemPtr) {  
    if (array[tail].load(memory_order_relaxed) != 0)  
        return false;  
    array[tail].store(itemPtr, memory_order_relaxed);  
    tail = (tail + 1) % size;  
    return true;  
}  
bool pop(Item * & itemPtr) {  
    itemPtr = array[head].load(memory_order_relaxed);  
    if (itemPtr == 0)  
        return false;  
    array[head].store(0, memory_order_relaxed);  
    head = (head + 1) % size;  
    return true;  
}
```


*SPSC Circular Queue of **non-null pointers***



```
std::atomic<Item *> array[size];
```

```
bool push(Item * itemPtr) {  
    if (array[tail].load(memory_order_relaxed) != 0)  
        return false;  
    array[tail].store(itemPtr, memory_order_relaxed);  
    tail = (tail + 1) % size;  
    return true;  
}  
bool pop(Item * & itemPtr) {  
    itemPtr = array[head].load(memory_order_relaxed);  
    if (itemPtr == 0)  
        return false;  
    array[head].store(0, memory_order_relaxed);  
    head = (head + 1) % size;  
    return true;  
}
```

SPSC Circular Queue of *non-null pointers*



```
std::atomic<Item *> array[size];

bool push(Item * itemPtr) {
    if (array[tail].load(memory_order_relaxed) != 0)
        return false;
    array[tail].store(itemPtr, memory_order_release);
    tail = (tail + 1) % size;
    return true;
}

bool pop(Item * & itemPtr) {
    itemPtr = array[head].load(memory_order_acquire);
    if (itemPtr == 0)
        return false;
    array[head].store(0, memory_order_relaxed);
    head = (head + 1) % size;
    return true;
}
```

Conclusions...

A Guide to Threaded Coding

“Lock-free coding is the last thing you want to do.”

1. Forget what you learned in Kindergarten
(stop Sharing)
2. Use Locks
3. Measure
4. Measure
5. Change your Algorithm
6. GOTO 1

- ∞. Lock-free

“Use Locks!”